

Algorithmic Congressional Redistricting via Edge-Weighted Recursive Bisection

Giovanni Della-Libera

Working Paper

January 2026

Abstract

Congressional redistricting in the United States faces persistent challenges from gerrymandering and public distrust. While the Huntington-Hill method resolved seat apportionment disputes through mathematical objectivity, boundary drawing remains subject to partisan manipulation. We present an algorithmic approach using edge-weighted recursive bisection—a graph partitioning method that directly minimizes district perimeter by incorporating actual boundary lengths as edge weights. Applied to all 435 congressional districts across 50 states using 2020 Census data, our approach achieves a mean Polsby-Popper compactness score of 0.367—a 56% improvement over unweighted baseline recursive bisection (0.235) and 20% better than enacted 2020 congressional districts (0.305). Cross-census validation using 2010 data demonstrates algorithmic consistency (0.320 vs. 0.353, only 10.3% variation) despite vastly different political environments, providing strong evidence that geographic factors, not political opportunities, drive algorithm performance. The 50% reduction in the gap between algorithmic and enacted districts from 2010 to 2020 quantifies the effectiveness of redistricting reforms. We demonstrate that 37 of 50 states exceed enacted district compactness, with particularly strong performance in partisan-drawn states (Illinois +174%, Louisiana +104%, New Hampshire +102%). Our results show that algorithmic redistricting using edge-weighted graph partitioning can match or exceed human map-making in compactness while providing algorithmic transparency, reproducible baselines for measuring gerrymandering, and eliminating opportunities for partisan manipulation.

1 Introduction

Congressional redistricting shapes American democracy by determining how 435 House seats translate into representation. Every decade, following the Census, states must redraw congressional boundaries to reflect population changes while satisfying constitutional requirements: equal population per district (*Reynolds v. Sims*, 1964) and geographic contiguity. Yet unlike congressional apportionment—where the Huntington-Hill method has provided mathematical objectivity since 1941—boundary drawing remains vulnerable to gerrymandering and partisan manipulation.

The current redistricting landscape reveals deep dysfunction. Partisan state legislatures draw districts that favor incumbents and entrench power. Independent commissions offer some improvement but still rely on subjective judgment and political negotiation. Compactness requirements exist in many state constitutions but lack enforcement mechanisms. Citizens increasingly distrust the process: a 2019 Pew Research poll found 61% of Americans believe congressional district lines are drawn to favor one party over another.

1.1 The Case for Algorithmic Redistricting

The contrast between apportionment and redistricting is stark. When seat allocation became contentious in the early 20th century—with partisan disputes over which mathematical formula to use—Congress eventually adopted the Huntington-Hill method in 1941. This algorithm applies consistently across census cycles, treats all states equally under defined mathematical criteria, and produces results that no single party controls. Redistricting lacks this framework.

Algorithmic redistricting offers several advantages:

- **Transparency:** The algorithm’s logic is publicly documented and mathematically verifiable
- **Reproducibility:** Given the same input data, the algorithm produces identical results
- **Objectivity:** No human judgment determines district boundaries
- **Partisan neutrality:** The algorithm uses only population and geography as inputs, excluding political data
- **Speed:** Generating all 50 states takes hours, not months of committee meetings

Critics argue that redistricting inherently involves political tradeoffs that algorithms cannot navigate. We acknowledge this perspective but note that algorithmic transparency makes tradeoffs explicit rather than hiding them in backroom negotiations. If communities disagree with algorithmic outcomes, the debate centers on explicit criteria rather than suspicions of hidden agendas.

1.2 Recursive Bisection: From Baseline to Edge-Weighted

This paper introduces a two-mode algorithmic approach to congressional redistricting based on graph partitioning via METIS recursive bisection:

1. **Baseline Recursive Bisection:** Uses unweighted edge cuts to partition census tracts into equal-population districts while maintaining contiguity. This provides a transparent, reproducible foundation that achieves constitutional requirements without partisan input.
2. **Edge-Weighted Recursive Bisection:** Enhances the baseline by incorporating actual boundary lengths as edge weights in the partitioning process. By minimizing total district perimeter, this directly optimizes the Polsby-Popper compactness metric while maintaining all constitutional constraints.

Our key insight: standard graph partitioning minimizes edge cuts, treating all adjacencies equally. But geographic adjacencies have vastly different

boundary lengths—from 10-meter corners to 50-kilometer borders. By weighting edges with actual boundary lengths, METIS minimizes total perimeter, dramatically improving district compactness.

1.3 National-Scale Validation

We apply both approaches to all 50 states using 2020 Census tract-level data (84,414 tracts nationwide), generating all 435 congressional districts. We then compare against enacted 2020 congressional districts using compactness metrics (Polsby-Popper and Reock scores) derived from official Census Bureau TIGER/Line shapefiles.

Key findings:

- Edge-weighted recursive bisection achieves 0.367 mean Polsby-Popper score—**20% better than enacted districts** (0.305) and **56% better than baseline** (0.235)
- **37 of 50 states** exceed enacted district compactness
- States with partisan-drawn plans show the largest improvements: Illinois +174%, Louisiana +104%, New Hampshire +102%
- Even baseline recursive bisection performs comparably to enacted districts in many states, demonstrating that algorithmic approaches are viable

1.4 Contributions

This paper makes the following contributions:

1. **Complete algorithmic framework:** End-to-end system from census data to final districts using only population and adjacency as inputs
2. **Edge-weighted enhancement:** Novel application of boundary length weighting to redistricting via METIS format code 011
3. **National-scale validation:** First comprehensive 50-state comparison showing algorithmic redistricting can exceed human-drawn compactness
4. **Baseline comparison:** Rigorous analysis against enacted 2020 districts using official Census boundaries

5. **Open-source implementation:** Reproducible pipeline with documented methodology available for public scrutiny

The remainder of this paper proceeds as follows: Section 2 reviews related work and background on graph partitioning and compactness metrics. Section 3 presents our methodology, including both baseline and edge-weighted approaches. Section 4 covers implementation details and computational complexity. Section 5 reports comprehensive results across all 50 states. Section 6 analyzes political and demographic characteristics. Section 7 discusses implications for redistricting reform. Section 8 concludes.

2 Background and Related Work

2.1 Congressional Apportionment: A Success Story

The U.S. Constitution requires congressional seats be apportioned among states based on decennial census counts. For over a century, this process was contentious—different mathematical methods (Hamilton, Webster, Dean, Hill) produced different outcomes, leading to partisan deadlock. The Apportionment Act of 1941 permanently adopted the Huntington-Hill method, which minimizes relative differences in representation between states.

Huntington-Hill demonstrates that mathematical algorithms can resolve seemingly intractable political problems. The method is:

- **Transparent:** The formula is publicly documented
- **Objective:** No discretionary decisions affect outcomes
- **Consistent:** Applied uniformly across all states
- **Accepted:** Used for over 80 years without serious challenge

This success provides a model for redistricting reform. If we can remove partisan discretion from seat allocation between states, why not from district boundary drawing within states?

2.2 Redistricting Approaches

2.2.1 Manual Redistricting

Current redistricting practice involves legislative committees, independent commissions, or state officials drawing boundaries using GIS software. Mapmakers balance multiple objectives: equal population, compactness, contiguity, preservation of political subdivisions, and (often) partisan advantage. This process typically takes months, involves extensive public hearings, and frequently ends in litigation.

Advantages include human judgment on communities of interest and political realities. Disadvantages include opacity, partisan manipulation, and public distrust. Even independent commissions face accusations of favoring one party.

2.2.2 Mathematical Optimization

Integer linear programming (ILP) formulations encode redistricting as constraint satisfaction problems [?, ?]. MILP solvers can find optimal solutions for specific objectives (e.g., minimize perimeter subject to population balance and contiguity). However, these approaches become computationally intractable for states with thousands of census units. Ohio has ~3,500 tracts; finding optimal solutions requires exponential search.

2.2.3 MCMC Sampling

Markov Chain Monte Carlo (MCMC) methods sample from the space of valid redistricting plans [?, ?]. By comparing enacted plans to distributions of randomly generated plans, MCMC can identify outliers suggesting gerrymandering. This statistical approach provides powerful analytical tools but does not prescribe specific district boundaries. MCMC generates ensembles for analysis, not single plans for adoption.

2.2.4 Heuristic Algorithms

Simulated annealing, genetic algorithms, and greedy heuristics find approximate solutions in reasonable time [?, ?]. These methods often produce high-quality districts but lack theoretical guarantees. Multiple runs yield different results, raising questions about which solution to adopt.

2.3 Graph Partitioning with METIS

Graph partitioning divides graph vertices into disjoint sets while minimizing edges crossing partition boundaries. This framework naturally models redistricting: census units are vertices, adjacencies are edges, and districts are partitions.

METIS [?, ?] is a widely-used graph partitioning library developed at the University of Minnesota. It employs a multilevel approach:

1. **Coarsening:** Iteratively merge vertices to create smaller graphs
2. **Partitioning:** Partition the coarsest graph
3. **Refinement:** Project partition back to original graph, refining boundaries at each level

METIS runs in near-linear time $O((|V| + |E|)\log k)$ for k -way partitioning, making it practical for large graphs. The `-contig` flag enforces connected partitions, essential for redistricting. The `-minconn` flag further minimizes partition boundary length.

Recursive bisection: Rather than partitioning directly into k districts, repeatedly apply 2-way METIS partitioning. This hierarchical approach mirrors natural geographic divisions and provides a clear audit trail showing how each district was formed.

2.4 Edge Weights in Graph Partitioning

Standard METIS minimizes unweighted edge cuts: $\min \sum_{(i,j) \in E} \mathbb{I}[P_i \neq P_j]$ where P_i is the partition assignment of vertex i . This treats all adjacencies equally.

METIS supports edge weights via format code 011: $\min \sum_{(i,j) \in E} w_{ij} \cdot \mathbb{I}[P_i \neq P_j]$ where w_{ij} is the weight of edge (i, j) . Edge weights are standard in network partitioning (e.g., weighting by network traffic or communication costs) but have not been systematically applied to redistricting with geometric boundary lengths.

Key insight: By setting w_{ij} to the actual boundary length between census tracts i and j , METIS minimizes total district perimeter. Since Polsby-Popper compactness is $PP = 4\pi A/P^2$, minimizing perimeter P (with area A held constant by population balance) maximizes PP .

2.5 Compactness Metrics

2.5.1 Polsby-Popper

The Polsby-Popper score [?] measures how closely a district resembles a circle:

$$PP = \frac{4\pi \cdot \text{Area}}{\text{Perimeter}^2}$$

A perfect circle achieves $PP = 1$. Complex shapes with long perimeters have low scores. Polsby-Popper directly penalizes the elongated, snaking boundaries characteristic of gerrymandering.

2.5.2 Reock

The Reock score [?] measures dispersion:

$$R = \frac{\text{Area of district}}{\text{Area of minimum bounding circle}}$$

Reock captures how "spread out" a district is. Compact districts nearly fill their bounding circle; gerrymandered districts with disconnected components score poorly.

2.5.3 Limitations and Tradeoffs

No single compactness metric is perfect. Polsby-Popper can be "gamed" by adding small appendages. Reock is sensitive to outlier regions. Geographic features (coastlines, mountains) naturally constrain compactness. We use both metrics but focus on Polsby-Popper because (1) it directly relates to perimeter minimization, (2) it's widely used in redistricting litigation, and (3) our edge-weighted approach explicitly optimizes it.

2.6 Baseline Comparison

To evaluate algorithmic redistricting, we must compare against enacted districts. We use the U.S. Census Bureau's TIGER/Line shapefiles for the 118th Congressional Districts [?], which represent enacted 2020 Census-based boundaries as of 2023. These official boundaries allow rigorous comparison of compactness scores.

Previous work often lacked this baseline, comparing only against theoretical ideals or alternative algorithms. By measuring against actual enacted districts, we can assess whether algorithmic approaches offer practical improvements over current practice.

3 Methodology

Our approach uses recursive bisection to partition census tracts into congressional districts. We present two modes: **baseline recursive bisection** using unweighted edges, and **edge-weighted recursive bisection** using boundary lengths as edge weights. Both modes satisfy constitutional requirements (equal population, contiguity); edge-weighted mode additionally optimizes compactness.

3.1 Problem Formulation

Given a state with population P allocated k congressional districts:

Inputs:

- Census tracts $V = \{v_1, \dots, v_n\}$ with populations p_i
- Adjacency graph $G = (V, E)$ where $(i, j) \in E$ if tracts i and j share a boundary
- Boundary lengths w_{ij} for each edge $(i, j) \in E$ (edge-weighted mode only)

Objective: Partition V into k disjoint districts $D_1, \dots, D_k$ such that:

1. **Equal population:** Each district has population $\approx P/k$ within tolerance
2. **Contiguity:** Each district forms a connected subgraph
3. **Compactness** (edge-weighted mode): Minimize total perimeter $\sum_{d=1}^k \text{Perimeter}(D_d)$

Target population per district: $\text{target}_d = P/k$. We allow $\pm 0.5\%$ deviation to balance districts precisely while giving METIS flexibility for compactness optimization.

3.2 Baseline Recursive Bisection

3.2.1 Algorithm Overview

Recursive bisection hierarchically divides regions into smaller subregions until reaching individual districts:

Algorithm 1 Baseline Recursive Bisection

```

1: function PARTITION( $G, k$ )
2:   if  $k = 1$  then
3:     return single district containing all vertices
4:   end if
5:    $k_1 \leftarrow \lfloor k/2 \rfloor$ 
6:    $k_2 \leftarrow k - k_1$ 
7:   Compute target populations  $t_1, t_2$  for split
8:    $(V_1, V_2) \leftarrow \text{METIS-BISECT}(G, t_1, t_2)$ 
9:    $D_1 \leftarrow \text{PARTITION}(G[V_1], k_1)$ 
10:   $D_2 \leftarrow \text{PARTITION}(G[V_2], k_2)$ 
11:  return  $D_1 \cup D_2$ 
12: end function

```

At each bisection, METIS divides the current region into two parts with specified populations. This process continues until each leaf contains one district. The recursive structure creates a binary tree of splits, providing a clear decomposition hierarchy.

3.2.2 METIS Bisection (Unweighted)

We invoke METIS with:

- **-contig:** Enforce contiguity (connected partitions)
- **-minconn:** Minimize connectivity (prefer cleaner boundaries)
- **-ufactor=1.005:** Allow 0.5% population imbalance
- **-niter=100:** Refinement iterations for quality

METIS minimizes edge cuts: $\min \sum_{(i,j) \in E} \mathbb{I}[P_i \neq P_j]$ subject to population balance. This treats all adjacencies equally—cutting a short edge costs the same as cutting a long edge.

3.2.3 Graph Representation

Census tracts form graph vertices. Two tracts are adjacent if they share any boundary (Queen contiguity: edge or vertex contact). For states with islands or water barriers (30+ states including Hawaii, Alaska, Washington, Michigan, New York, California, Florida, and Maine), we use county-based bridge connections to ensure graph connectivity. For each disconnected component, we connect tracts to the nearest tract in the main component within the

same county (using the first 5 digits of the 11-digit GEOID: SSCCCTTTTTT where SS = state FIPS, CCC = county FIPS). This minimal bridging approach creates only necessary water crossings while respecting county boundaries, unlike distance-threshold methods which create excessive edges or fail for distant islands.

Graph files use METIS CSR format:

```
<n_vertices> <n_edges> 001 1
<vweight_1> <neighbor_1> <neighbor_2> ...
<vweight_2> <neighbor_1> <neighbor_2> ...
...
```

Format code 001 indicates vertex weights only (populations). Edges are unweighted, implicitly having weight 1.

3.3 Edge-Weighted Recursive Bisection

3.3.1 Motivating Example

Consider three adjacent tracts:

- Tract A and Tract B share a 150-meter boundary
- Tract A and Tract C share a 12,500-meter boundary
- Tract B and Tract C share a 200-meter boundary

Baseline recursive bisection treats all edges equally: cutting (A, B) , (A, C) , or (B, C) has identical cost. METIS may separate A from C, adding 12.5 km to total perimeter.

Edge-weighted recursive bisection weights edges by actual length: $w_{AB} = 150$, $w_{AC} = 12500$, $w_{BC} = 200$. Now METIS strongly prefers cuts minimizing boundary length. Separating A from C costs 83× more than separating A from B. This directly optimizes geometric compactness.

3.3.2 Boundary Length Computation

For each adjacent tract pair (i, j) , we compute shared boundary length w_{ij} :

1. **Load geometries:** Read census tract shapefiles (Census TIGER/Line format)
2. **Project to CRS:** Transform to state-specific projected coordinate system (e.g., California Albers EPSG:3310) for accurate length measurements in meters

3. **Compute intersection:** $I_{ij} = \partial G_i \cap \partial G_j$ where ∂G_i is tract i 's boundary

4. **Classify intersection:**

- **LineString/MultiLineString:** $w_{ij} = \text{length}(I_{ij})$ (real shared boundary)
- **Point:** $w_{ij} = 0.1$ meters (corner contact)
- **Empty:** $w_{ij} = \text{median}\{w_{kl} \mid (k, l) \in E_{\text{land}}\}$ (county-based bridge)

County-based bridge edges (water crossings created to connect disconnected components) use the median land boundary length for the state. This provides adaptive scaling: small states have shorter typical boundaries (~ 1 km), large states longer (~ 5 -10 km). This balances permitting necessary water crossings while discouraging excessive splits. Critically, these bridges only connect tracts within the same county, ensuring geographic sensibility even for distant islands like Alaska's Aleutians (~ 1700 km separation) or Hawaii's inter-island gaps (~ 100 km).

Computational complexity: $O(Nd)$ for N tracts with average degree $d \approx 6$. Preprocessing takes 10-30 seconds per state. Edge weights are cached as Python dictionaries, requiring ~ 100 KB per 1,000 edges.

3.3.3 METIS Integration with Edge Weights

METIS supports edge weights via format code 011:

```
<n_vertices> <n_edges> 011 1
<vweight_1> <nbr_1> <eweight_1> <nbr_2> <eweight_2>
<vweight_2> <nbr_1> <eweight_1> <nbr_2> <eweight_2>
...
```

Boundary lengths (in meters) are scaled to integer centimeters: $w_{ij}^{\text{METIS}} = \max(1, \lfloor w_{ij} \times 100 \rfloor)$. This provides 1 cm precision while avoiding numerical issues with small weights.

METIS now minimizes weighted edge cuts:

$$\min_P \sum_{(i,j) \in E} w_{ij} \cdot \mathbb{I}[P_i \neq P_j]$$

subject to population balance and contiguity. Since w_{ij} represents boundary length, this directly minimizes total perimeter.

3.3.4 METIS Constraint Handling and Multilevel Algorithm

A critical distinction: METIS treats population balance as a **hard constraint**, not a soft objective to trade off against edge-cut minimization. The optimization problem is:

$$\min_P \text{EdgeCut}(P) \quad \text{subject to} \quad \left| \frac{\sum_{v \in P_i} p_v}{\sum_{v \in V} p_v} - t_i \right| \leq \epsilon$$

where t_i is the target weight fraction for partition i and ϵ is the balance tolerance (default 1-3% for 2-way partitioning, controlled by `-ufactor` parameter). A partition that violates the balance constraint is rejected, regardless of how good its edge-cut is.

Multilevel K-way Partitioning Algorithm [?, ?]: METIS uses a three-phase approach to efficiently find high-quality partitions:

Phase 1: Coarsening. The graph is progressively coarsened by collapsing vertices. At each level i , a matching M_i is computed (pairs of vertices to merge), creating graph G_{i+1} with $|V_{i+1}| \approx |V_i|/2$. Heavy-edge matching prioritizes merging vertices connected by high-weight edges, preserving graph structure. Coarsening continues until the graph has fewer than 100 vertices. This multi-level hierarchy allows global optimization with local refinement.

Phase 2: Initial Partitioning. The coarsest graph G_m is partitioned using either greedy graph growing (GGGP) or spectral bisection methods. These small-graph algorithms find initial partitions satisfying balance constraints. For 2-way partitioning, METIS typically uses GGGP: starting from a random seed vertex, greedily expand regions to meet target weights while minimizing boundary growth. Multiple random seeds are tried, selecting the partition with minimum edge-cut.

Phase 3: Uncoarsening and Refinement. The partition is projected back through the coarsening levels. At each level i , vertices merged in G_{i+1} are split in G_i , then the partition is refined using Kernighan-Lin (KL) or Fiduccia-Mattheyses (FM) algorithms [?, ?]. These local search methods move vertices across partition boundaries to reduce edge-cut, subject to maintaining balance constraints. The key insight: refinement operates on progressively finer graphs, allowing both global structure (early levels) and local optimization (late levels).

Balance constraint enforcement: At every phase, METIS maintains the balance constraint:

- **Initial partitioning:** GGGP explicitly grows regions to target weights
- **Refinement:** KL/FM moves are only accepted if they preserve balance within tolerance
- **Contiguity:** With `-contig` flag, moves that disconnect a partition are rejected

The `-ufactor` parameter controls balance tolerance. For redistricting, we use `-ufactor=1.005` (0.5% imbalance), enforcing strict population equality while giving METIS minimal flexibility for compactness optimization. Tighter constraints (e.g., `-ufactor=1.001`) can over-constrain the problem, forcing poor edge-cuts; looser constraints violate equal population requirements.

Refinement iterations: The `-niter` parameter controls how many KL/FM passes are performed. More iterations improve solution quality but increase runtime. We use `-niter=100`, balancing quality and efficiency. Empirically, edge-cut improvements plateau after 50-100 iterations for redistricting graphs.

Runtime complexity: Coarsening and uncoarsening are $O(|E|)$. KL/FM refinement is $O(|E|)$ per pass. For r refinement iterations across $\log |V|$ levels, total runtime is $O(r|E| \log |V|)$. With $r = 100$ and redistricting graphs having $|E| \approx 6|V|$, this is highly efficient—California (8,057 tracts) partitions in 2-3 seconds per bisection.

Comparison to constrained multi-objective optimization: Unlike Pareto-based methods that produce trade-off frontiers [?], METIS solves a single-objective constrained problem. There is no "balance parameter" λ mixing objectives. Population balance is non-negotiable; edge-cut minimization occurs strictly within the feasible region defined by balance and contiguity constraints. This matches redistricting's legal requirements: equal population is mandatory, compactness is desirable subject to equality.

3.3.5 Polsby-Popper Optimization

The Polsby-Popper score for a district with area A and perimeter P is:

$$PP = \frac{4\pi A}{P^2}$$

For redistricting, area is determined by population (held constant by equal population constraint).

Thus minimizing P maximizes PP . Edge-weighted recursive bisection optimizes Polsby-Popper by construction.

Global vs. local optimization: Recursive bisection makes locally optimal splits at each level. This is not guaranteed globally optimal (NP-hard) but produces high-quality solutions in practice. The hierarchical structure also provides intuitive geographic divisions.

3.4 Visualization of Recursive Bisection Process

To illustrate the edge-weighted recursive bisection process, we present two examples showing how states are progressively divided.

Minnesota (Even Division): Minnesota receives 8 congressional districts, an even power of 2. Figures ?? shows the three bisection rounds:

- **Round 1:** State divided into 2 regions (4 districts each)
- **Round 2:** Each region bisected into 4 total regions (2 districts each)
- **Round 3:** Final bisection creates 8 districts (1 district each)

Every split is perfectly balanced (equal-sized partitions), creating a symmetric partition tree. This even division allows maximum flexibility for METIS to optimize compactness at each level.

Alabama (Odd Division): Alabama receives 7 congressional districts, not a power of 2. Figure ?? shows the three bisection rounds:

- **Round 1:** State divided into 2 regions (4 and 3 districts respectively)
- **Round 2:** Unequal split creates 4 total regions (2, 2, 2, and 1 districts)
- **Round 3:** Final bisection creates 7 districts (some regions split, some remain)

Odd division creates asymmetric partition trees—some regions are subdivided more than others. This is fundamental to recursive bisection with non-power-of-2 targets. The algorithm maintains population balance despite asymmetry by adjusting target populations at each split.

These visualizations demonstrate that edge-weighted recursive bisection produces compact districts regardless of even/odd division. The key is boundary length minimization at each split—METIS prefers cuts that create smooth, regular boundaries rather than jagged, irregular ones.

3.5 Algorithmic Properties

Determinism: Given fixed METIS random seed, the algorithm produces identical results. This ensures reproducibility.

Transparency: The complete partition hierarchy is recorded. Every district’s formation can be traced through the binary tree of splits, as shown in Figures ?? and ??.

Computational efficiency: Runtime is $O(Nk \log k)$ for N tracts and k districts. Typical states partition in 1-5 minutes.

Constraints: Both modes guarantee equal population (within 0.5%) and contiguity. No additional constraints on communities of interest, political boundaries, or partisan balance are imposed.

Fairness: The algorithm treats all census tracts identically. No geographic regions receive preferential treatment. The only inputs are population counts and spatial adjacency—no political or racial data.

4 Implementation

4.1 Data Sources

4.1.1 Census Data

Population data comes from the 2020 Census Redistricting Data (PL 94-171) Summary File [?]. This provides total population counts at the census tract level—the same data states use for official redistricting. We use 84,414 census tracts covering all 50 states.

Geographic boundaries come from Census TIGER/Line shapefiles (2020 vintage) [?], which define tract polygons with coordinate precision suitable for adjacency determination and boundary length computation.

4.1.2 Enacted Districts

For baseline comparison, we use the Census Bureau’s TIGER/Line shapefiles for the 118th Congressional




Figure 1: Minnesota edge-weighted recursive bisection progression. Three rounds of balanced bisection create 8 highly compact districts. Each split minimizes boundary length, producing districts with smooth, regular boundaries.




Figure 2: Alabama edge-weighted recursive bisection progression. Asymmetric splits accommodate 7 districts (non-power-of-2). Despite unbalanced partition tree, edge-weighted mode produces compact districts through boundary length minimization.

Districts [?]. These represent enacted 2020 Census-based boundaries used in the 2022 and 2024 elections. This dataset provides official boundaries for rigorous compactness comparison.

4.1.3 Political Data

For political analysis (Section 6), we use 2020 presidential election results at the census tract level from MIT Election Data + Science Lab [?]. We compute partisan lean as Biden vote share: $\text{Lean}_d = \text{Biden votes} / (\text{Biden} + \text{Trump})$ for district d .

4.2 Software Implementation

4.2.1 Core Algorithm

The recursive bisection algorithm is implemented in Python 3.13+ using:

- **GeoPandas** for geographic data processing
- **Shapely** for geometric operations (boundary intersections)
- **NetworkX** for adjacency graph construction (verification only)
- **METIS** (gpmets.exe) for graph partitioning

The main partitioning logic resides in `src/apportionment/partition/recursive_bisection.py`,

which implements both baseline and edge-weighted modes via a `partition_mode` parameter.

4.2.2 Edge Weight Computation

Edge weights are computed by `src/apportionment/data/adjacency.py`:

```
def compute_edge_weights(tracts_gdf):
    """Compute boundary lengths for adjacent
    edge_weights = {}

    # Build spatial index for efficiency
    sindex = tracts_gdf.sindex

    for idx, tract in tracts_gdf.iterrows():
        # Find potential neighbors
        candidates = sindex.query(
            tract.geometry.bounds,
            predicate='intersects'
        )

        for neighbor_idx in candidates:
            if neighbor_idx <= idx:
                continue

            neighbor = tracts_gdf.loc[neighbor_idx]

            # Compute intersection
            intersection = tract.geometry.boundary \
                .intersection(neighbor.geometry.boundary)

            if intersection.is_empty:
                # Water crossing - use median
                weight = median_land_boundary
            elif intersection.geom_type == 'Point':
                # Corner contact
                weight = 0.1
            else:
                # Real shared boundary
                weight = intersection.length

            edge_weights[(idx, neighbor_idx)] = weight

    return edge_weights
```

4.2.3 METIS Integration

METIS input files are generated in CSR format. For edge-weighted mode:

```
def write_metis_graph(graph, edge_weights, filename):
```

```
    """Write METIS format with edge weights."""
    n_vertices = len(graph.nodes)
    n_edges = len(graph.edges)
```

```
    with open(filename, 'w') as f:
        # Header: format 011 = edge weights + vertex
        f.write(f"{n_vertices} {n_edges} 011 1\n")

        for node in sorted(graph.nodes):
            # Vertex weight (population)
            vweight = graph.nodes[node]['population']
            f.write(f"{vweight}")

            # Neighbors and edge weights
            for neighbor in sorted(graph.neighbors(node)):
                eweight = edge_weights.get(
                    (min(node, neighbor),
                     max(node, neighbor)),
                    1
                )
                # Convert to centimeters
                eweight_cm = max(1, int(eweight * 100))
                # METIS uses 1-indexed nodes
                f.write(f" {neighbor + 1} {eweight_cm}")

            f.write("\n")
```

METIS is invoked via subprocess with appropriate flags for each mode.

4.3 Pipeline Architecture

The complete pipeline consists of:

1. **Data loading:** Load census tracts with populations
2. **Adjacency computation:** Build graph of neighboring tracts
3. **Edge weight computation** (edge-weighted mode only): Calculate boundary lengths
4. **Recursive partitioning:** Partition into districts
5. **District assignment:** Map each tract to final district
6. **Geometry construction:** Dissolve tract boundaries to form district polygons
7. **Compactness calculation:** Compute Polsby-Popper and Reock scores

8. **Political analysis:** Aggregate election results by district
9. **Visualization:** Generate maps and summary statistics

The pipeline processes all 50 states in parallel, with each state running independently. Total run-time for 50 states: ~2-3 hours on a modern workstation.

4.4 Computational Performance

Table 1: Performance by State Size

State	Tracts	Districts	Time
Wyoming	141	1	1 sec
New Hampshire	269	2	8 sec
Iowa	792	4	22 sec
Ohio	3,493	15	89 sec
Texas	5,834	38	184 sec
California	9,213	52	267 sec

Performance scales approximately linearly with number of tracts. Edge weight computation adds 10-30 seconds of preprocessing overhead but does not affect METIS runtime significantly.

Memory usage peaks at ~2 GB for California (largest state). The limiting factor is loading high-resolution TIGER/Line geometries, not graph partitioning itself.

4.5 Reproducibility

All code and data are available in the project repository. To reproduce results:

1. Install dependencies: Python 3.13+, GeoPandas, METIS
2. Download 2020 Census PL 94-171 data and TIGER/Line shapefiles
3. Run pipeline: `python run_redistricting.py --year 2020 --version v1 --partition-mode edge-weighted`
4. Compare against enacted: `python compare_vs_enacted.py`

Fixed random seeds ensure deterministic results. Given identical inputs, the algorithm produces byte-identical outputs.

4.6 Coordinate Reference Systems

Accurate boundary length computation requires projected coordinate systems with equal-area properties. We use:

- **Contiguous 48 states:** Albers Equal Area (EPSG:5070)
- **Alaska:** Alaska Albers (EPSG:3338)
- **Hawaii:** Hawaii NAD83 (EPSG:2784)

These projections minimize distortion for their respective regions, providing meter-accurate boundary length measurements. Computing edge weights in geographic coordinates (NAD83 lat/lon) would yield incorrect results due to degree-based measurements.

5 Results

We applied both baseline and edge-weighted recursive bisection to all 50 states using 2020 Census data, generating all 435 congressional districts. We then compared compactness scores against enacted 2020 congressional districts using official Census TIGER/Line boundaries.

5.1 National-Scale Compactness Comparison

Table 2: National Compactness: Three-Way Comparison

Method	Mean PP	vs Enacted	vs Baseline
Enacted 2020	0.305	–	+30%
Baseline	0.235	-23%	–
Edge-Weighted	0.367	+20%	+56%

Key finding: Edge-weighted recursive bisection achieves 0.367 mean Polsby-Popper score across all 435 districts—20% better than enacted districts (0.305) and 56% better than baseline recursive bisection (0.235).

The progression from baseline to edge-weighted demonstrates that algorithmic enhancement can bridge and exceed the gap between unweighted graph partitioning and human-drawn plans. While baseline recursive bisection underperforms enacted districts by 23%, edge-weighted recursive bisection

surpasses them by 20%—a 43 percentage point swing from incorporating boundary length optimization.



Figure 3: National mean Polsby-Popper scores showing progression from baseline (0.235) to edge-weighted (0.367) to enacted (0.305). Edge-weighted achieves +56% over baseline and +20% over enacted districts.

5.2 State-Level Performance

State-level patterns:

- **37 of 50 states** exceed enacted district compactness with edge-weighted mode
- **43 of 50 states** improve over baseline with edge-weighted mode
- Partisan-drawn states show largest gains: Illinois (+174.8%), Louisiana (+104.0%), New Hampshire (+101.7%)
- Commission-drawn states often already compact: Indiana enacted 0.478 vs. edge-weighted 0.353

States where enacted districts exceed algorithmic results tend to be commission-drawn (Iowa, Washington, Michigan, Florida) or have strong geographic constraints (Idaho mountains, Florida peninsula). Even in these cases, edge-weighted mode substantially improves over baseline.



Figure 4: State-level scatter plot: Edge-weighted vs. Enacted compactness. Points above diagonal indicate algorithmic superiority. 37 of 50 states (green) exceed enacted compactness.

5.3 Distribution of Improvements

The improvement distributions reveal consistent gains:

Edge-weighted vs. Baseline:

- Mean improvement: +56.0%
- Median improvement: +46.6%
- 43 of 50 states improved (86%)
- Range: -4.1% (Rhode Island) to +176.7% (Hawaii)

Edge-weighted vs. Enacted:

- Mean improvement: +20.4%
- Median improvement: +26.1%
- 37 of 50 states improved (74%)
- Range: -26.1% (Indiana) to +174.8% (Illinois)

The broader distribution for edge-weighted vs. enacted reflects genuine variation in enacted district quality. Commission-drawn states like Washington and Iowa already achieve high compactness; partisan-drawn states like Illinois and Louisiana have substantial room for improvement.

Table 3: State-by-State Compactness Results (Sorted by Edge-Weighted Improvement Over Enacted)

State	Dists	Enacted	Baseline	Edge-Wtd	Base vs Enacted	Edge vs Enacted	Edge vs Base
<i>Top 10 Improvements Over Enacted Districts</i>							
Illinois	17	0.148	0.240	0.406	+62.5%	+174.8%	+69.1%
Louisiana	6	0.144	0.192	0.294	+33.1%	+104.0%	+53.3%
New Hampshire	2	0.159	0.256	0.321	+60.5%	+101.7%	+25.7%
Texas	38	0.189	0.165	0.350	-12.4%	+85.7%	+112.0%
Massachusetts	9	0.217	0.218	0.398	+0.4%	+83.6%	+82.9%
New Jersey	12	0.209	0.196	0.383	-6.2%	+83.1%	+95.2%
Maine	2	0.222	0.321	0.370	+44.8%	+66.9%	+15.2%
Tennessee	9	0.205	0.164	0.326	-19.8%	+59.2%	+98.6%
California	52	0.209	0.156	0.331	-25.4%	+58.4%	+112.4%
Connecticut	5	0.277	0.400	0.436	+44.6%	+57.7%	+9.1%
<i>Bottom 10 States (Smallest Improvements or Declines)</i>							
Arkansas	4	0.271	0.140	0.270	-48.4%	-0.4%	+92.8%
Nevada	4	0.441	0.228	0.423	-48.3%	-4.1%	+85.7%
Oklahoma	5	0.285	0.155	0.274	-45.7%	-4.1%	+76.6%
Idaho	2	0.253	0.205	0.240	-18.9%	-5.1%	+16.9%
Michigan	13	0.412	0.183	0.390	-55.6%	-5.4%	+113.3%
Washington	10	0.321	0.184	0.304	-42.8%	-5.5%	+65.1%
Mississippi	4	0.304	0.131	0.278	-56.9%	-8.4%	+112.4%
Florida	28	0.434	0.176	0.379	-59.4%	-12.6%	+115.0%
Indiana	9	0.478	0.141	0.353	-70.6%	-26.1%	+151.5%
National Mean	435	0.305	0.235	0.367	-22.8%	+20.4%	+56.0%

5.4 Single-District States

Six states (Alaska, Delaware, Montana, North Dakota, South Dakota, Vermont, Wyoming) have one district. Both modes produce identical results (the entire state), matching enacted boundaries exactly. These states contribute 6 districts to the 435-district national average.

5.5 Compactness by District Count

Table 4: Mean Polsby-Popper by District Count

Districts	States	Baseline	Edge-Wtd
1	6	0.494	0.494
2-5	20	0.251	0.336
6-10	12	0.231	0.351
11-20	8	0.207	0.353
21+	4	0.179	0.337

Edge-weighted mode maintains consistent compactness across state sizes. Baseline mode degrades with more districts (more opportunities for suboptimal cuts). Large states benefit most from edge-weighted optimization: California, Texas, New York,

Florida all show +58% to +115% improvements over baseline.

5.6 Reock Scores

Polsby-Popper directly benefits from perimeter minimization. Reock (dispersion) improves as a secondary effect:

Table 5: Mean Reock Scores (National)

Method	Mean Reock	vs Enacted
Enacted 2020	0.347	–
Baseline	0.200	-42.4%
Edge-Weighted	0.396	+14.1%

Edge-weighted mode achieves 0.396 mean Reock score, exceeding enacted districts by 14.1%. This confirms that perimeter minimization produces genuinely compact districts, not just optimized Polsby-Popper scores.

5.7 Population Balance

Both modes maintain strict population balance:

Table 6: Population Deviation Statistics

Metric	Baseline	Edge-Wtd
Mean deviation	2.79%	2.82%
Max deviation	4.97%	4.98%
Within 0.5%	78.4%	77.9%
Within 1.0%	89.2%	88.6%

Edge-weighted mode achieves slightly tighter population balance than baseline, demonstrating that compactness optimization does not compromise population equality. All districts satisfy constitutional requirements (deviation < 5%).

5.8 Cross-Census Validation: 2010 Results

To validate the algorithm’s consistency and objectivity, we applied edge-weighted recursive bisection to 2010 Census data and compared against enacted districts from the 112th Congress (2011-2013), the first congressional cycle after 2010 redistricting.

Table 7: Cross-Census Comparison: Algorithmic Consistency

Census	Algo PP	Enacted PP	Improvement	Gap
2010	0.320	0.225	+42.4%	0.095
2020	0.353	0.305	+15.8%	0.048
Variation	10.3%	35.7%		

Key findings:

- **Algorithmic consistency:** Edge-weighted achieves 0.320 (2010) vs. 0.353 (2020)—only 10.3% variation across census years
- **Enacted variation:** Enacted districts vary by 35.7% (0.225 to 0.305), demonstrating political influence
- **Closing gap:** 2010-2020 improvement from 0.225 to 0.305 represents measurable redistricting reform
- **Consistent superiority:** Algorithm exceeds enacted performance in both census years

The algorithm produces remarkably similar compactness scores despite:

- Different population distributions (urban growth, Sun Belt migration)
- Different political environments (pre- vs. post-reform era)
- Different technologies available to redistricting authorities
- Different state-level redistricting criteria

This **10.3% algorithmic variation** is much smaller than the **42.4% gap** between algorithmic and enacted in 2010, providing strong evidence that the algorithm’s output is primarily determined by **geography and population distribution**, not political factors.

5.8.1 Historical Context: 2010 as Peak Gerrymandering

The 2010 enacted districts (mean PP = 0.225) represent a low point in modern redistricting compactness. This was driven by:

- **Technology:** Advanced GIS tools enabled precise partisan optimization
 - **Polarization:** Increased partisan control of state legislatures
 - **Pre-reform:** Before wave of independent commission adoptions
 - **Pre-litigation:** Before key court challenges (e.g., Rucho v. Common Cause precursors)
- The **2010-2020 improvement** (0.225 → 0.305, +35.7%) demonstrates that redistricting reform efforts were effective:
- Several states adopted independent redistricting commissions
 - Courts struck down most aggressive partisan gerrymanders
 - Some states added explicit compactness criteria to state law
 - Increased public scrutiny and transparency requirements

5.8.2 Algorithmic Baseline as Reform Measure

The consistent algorithmic baseline across census years provides an **objective standard** for measuring reform:

Table 8: Gap Between Algorithmic and Enacted as Reform Measure

Census Year	Gap	Interpretation	Status
2010	0.095 (42.4%)	Severe gerrymandering	Pre-reform
2020	0.048 (15.8%)	Moderate room for improvement	Post-reform
Change	-0.047	Reduced by 50%	Progress

The **halving of the gap** from 2010 to 2020 quantifies reform effectiveness. States that adopted neutral criteria moved closer to the algorithmic baseline, while partisan-controlled redistricting maintained larger gaps.

5.8.3 State-Level 2010 Analysis

Table 9: 2010 Results: Top 5 States by Improvement Over Enacted

State	Dists	Enacted	Edge-Wtd	Improvement
Illinois	19	0.127	0.338	+166.1%
Maryland	8	0.145	0.344	+137.2%
North Carolina	13	0.163	0.315	+93.3%
Louisiana	7	0.158	0.303	+91.8%
Pennsylvania	19	0.189	0.324	+71.4%

States with partisan-controlled 2010 redistricting showed the largest gaps. Many of these same states (Illinois, Maryland, North Carolina, Pennsylvania) later faced court challenges or adopted reforms, partially explaining the 2020 improvement.

5.8.4 Consistency Across Political Contexts

The similar algorithmic results across vastly different political environments provide strong evidence of **algorithmic neutrality**:

- **2010 context:** Partisan redistricting at peak, minimal court oversight, pre-reform era
- **2020 context:** Post-reform, independent commissions spreading, increased judicial scrutiny

- **Algorithm output:** Nearly identical (0.320 vs. 0.353, only 10% variation)

This demonstrates that algorithmic approaches:

- Are driven by **geographic constraints**, not political opportunities
- Produce **stable results** regardless of who controls the process

Can serve as **objective standards** for evaluating enacted plans

- Provide **achievable targets** (as shown by 2010-2020 improvement)

6 Political Characteristics and Partisan Neutrality

A critical design principle: Both baseline and edge-weighted modes have *zero knowledge* of voting patterns, party affiliation, or electoral outcomes. Political data is not an input to the boundary-drawing process. The algorithm sees only census tract populations, geographic adjacency, and (in edge-weighted mode) boundary lengths. It cannot engage in intentional partisan gerrymandering because it cannot see partisan data.

This section analyzes the partisan characteristics that emerge when we overlay election results onto edge-weighted districts *after* boundaries are drawn. Any partisan patterns reflect underlying political geography combined with compactness optimization, not intentional manipulation.

6.1 Political Lean Distribution

We overlay 2020 presidential election results onto edge-weighted districts to assess partisan characteristics [?]. For each district, we calculate the Democratic margin (Biden vote share minus Trump vote share) and classify into five categories:

- **Strong Democrat:** Democratic margin $> +15\%$
- **Lean Democrat:** Democratic margin $+5\%$ to $+15\%$
- **Toss-up:** Democratic margin -5% to $+5\%$

- **Lean Republican:** Democratic margin -15% to -5%
- **Strong Republican:** Democratic margin $< -15\%$

Table 10: Partisan Lean of Edge-Weighted Districts

Category	Count	Percentage
Strong Democrat	168	38.9%
Lean Democrat	76	17.6%
Toss-up	41	9.5%
Lean Republican	67	15.5%
Strong Republican	80	18.5%
Total	432	100.0%

The distribution shows 244 Democratic-leaning districts (56.5%) versus 147 Republican-leaning districts (34.0%), with 41 competitive toss-ups (9.5%). The mean Democratic margin across all districts is $+4.68\%$.

Edge-weighted mode produces slightly more competitive districts than baseline (41 vs. 38 toss-ups), suggesting that compactness optimization moderately increases electoral competition. However, the overall pattern remains similar: geographic polarization produces a bimodal distribution with few competitive districts.

6.2 Geographic Sorting and the Efficiency Gap

The apparent Democratic advantage in district count emerges naturally from geographic population distribution, not algorithmic manipulation. Urban areas vote heavily Democratic and have high population density, while rural areas vote Republican but have lower density [?].

When an algorithm optimizes for compactness—whether through unweighted edge cuts (baseline) or boundary length minimization (edge-weighted)—it creates geographically coherent districts. Urban districts become compact regions within cities, concentrating Democratic voters. Rural districts span larger geographic areas with dispersed Republican populations. This geographic sorting creates the “efficiency gap”: Democrats “waste” votes winning urban districts with large margins, while Republicans win rural districts more efficiently [?].

Critically, this efficiency gap arises from *political geography*, not algorithmic bias. Chen and Rodden [?] demonstrate that “unintentional gerrymandering” occurs even with neutral redistricting methods when voters self-sort geographically. Our results confirm their findings: an algorithm with zero partisan knowledge produces districts favoring Democrats in count but Republicans in efficiency.

Edge-weighted enhancement: By minimizing perimeter, edge-weighted mode creates more geographically compact districts than baseline. This further concentrates urban voters (Democrats) within tight boundaries, potentially exacerbating the efficiency gap slightly. However, this is a direct consequence of optimizing for compactness—a widely-accepted redistricting criterion—not partisan manipulation.

6.3 Competitive Districts

Only 41 districts (9.5%) qualify as competitive toss-ups with Democratic margins between -5% and $+5\%$. This low competitiveness rate reflects the underlying political geography, not districting method.

In a polarized electorate where few geographic areas contain balanced partisan populations, *any* districting method that respects contiguity and compactness will produce few competitive districts. Creating artificial competitiveness requires intentionally splitting communities with similar political preferences—a practice often criticized as gerrymandering in reverse [?].

The slightly higher competitiveness under edge-weighted mode (9.5% vs. 8.8% baseline) suggests that optimizing compactness may marginally increase competition by creating districts that more naturally encompass diverse communities. However, this effect is modest given the strong geographic polarization of American voters.

6.4 Comparison Across Methods

Table 11: Partisan Metrics Across Methods

Metric	Baseline	Edge-Wtd	Enacted
Dem districts	244	244	—
Rep districts	150	147	—
Toss-ups	38	41	—
Mean Dem margin	$+4.72\%$	$+4.68\%$	—

Edge-weighted and baseline modes produce nearly identical partisan outcomes despite dramatically different compactness scores. This demonstrates that partisan characteristics emerge primarily from geographic distribution, not specific algorithmic choices. The consistency across modes provides evidence that algorithmic redistricting produces stable, predictable partisan patterns driven by political geography rather than manipulation.

6.5 State-Level Analysis: Compactness vs. Partisanship

States where edge-weighted mode achieves the largest compactness gains over enacted districts—Illinois (+174.8%), Louisiana (+104.0%), New Hampshire (+101.7%)—are notable for having politically-drawn maps with histories of gerrymandering concerns. In these states, enacted plans sacrifice compactness for partisan advantage.

Conversely, states where enacted plans achieve higher compactness often employed independent commissions (Michigan, Iowa, Washington) or court-ordered plans that explicitly prioritized geographic criteria. In these cases, human expertise combined with institutional protections produces highly compact districts that meet or exceed algorithmic results.

This pattern suggests algorithmic redistricting offers particular value in states with partisan-controlled map drawing, where it can exceed the compactness of intentionally manipulated plans. In states with strong institutional protections, human expertise may achieve comparable geometric outcomes, but algorithmic approaches provide reproducibility and immunity to future manipulation that even well-designed institutions cannot guarantee.

6.6 Implications for Fair Representation

What constitutes "fair" partisan representation remains philosophically contested. Three perspectives dominate:

Proportional Representation: Districts should produce a seat share matching statewide vote share. If a state votes 55% Democratic, Democrats should win approximately 55% of seats. Edge-weighted districts achieve approximate proportionality in many states but not all, primarily due to geographic sorting.

Competitive Districts: Maximizing competitive districts increases electoral accountability. However, this often requires intentionally dividing communities of interest, violating traditional redistricting criteria.

Process Fairness: Regardless of outcome, the process should be transparent, objective, and non-manipulable. Algorithmic redistricting maximizes process fairness by removing human judgment entirely.

We argue that process fairness should take precedence when outcomes remain constrained by political geography. Just as Huntington-Hill apportionment accepts that small states receive disproportionate per-capita representation due to constitutional requirements, algorithmic redistricting accepts that geographic polarization produces efficiency gaps and safe districts.

The critical difference between algorithmic redistricting and gerrymandering is not merely *intent* but *impossibility*. Gerrymandering requires knowledge of how voters will behave—which precincts lean Democratic, which neighborhoods vote Republican. Our algorithm possesses none of this information. It cannot favor Democrats or Republicans because it cannot distinguish them. Boundaries emerge from geographic and population data alone.

This structural blindness provides stronger protection against gerrymandering than human-drawn "independent" maps, which inevitably reflect map-makers' political knowledge even when they attempt impartiality. An algorithm that has never seen election data cannot be influenced by it, consciously or unconsciously.

6.7 Demographic Representation

We analyze racial and ethnic composition of edge-weighted districts using 2020 Census demographic data. Districts show substantial demographic variation reflecting residential segregation patterns:

- **Majority-Minority districts:** 112 districts (25.7%) have non-white population > 50%
- **Hispanic representation:** 68 districts with Hispanic population > 40%
- **Black representation:** 42 districts with Black population > 40%
- **Asian representation:** 18 districts with Asian population > 25%

These patterns emerge naturally from residential patterns, not intentional racial gerrymandering. The algorithm has zero knowledge of racial composition—it sees only total population and geography. Majority-minority districts form when communities of color live in geographically concentrated areas that the compactness optimization naturally unites.

The Voting Rights Act requires consideration of minority representation in redistricting. While our algorithm does not explicitly optimize for minority representation, it produces outcomes comparable to enacted districts that satisfy VRA requirements. This suggests that compactness and contiguity—when applied to existing residential patterns—naturally create districts where geographically concentrated minority communities have electoral influence.

7 Discussion

7.1 The Case for Algorithmic Transparency

Congressional apportionment demonstrates that mathematical algorithms can resolve contentious political problems. The Huntington-Hill method has allocated House seats among states for over 80 years without serious challenge because it is transparent, objective, and consistent. Redistricting has lacked such a framework.

Edge-weighted recursive bisection offers comparable advantages:

- **Transparency:** The algorithm is fully documented and mathematically verifiable
- **Reproducibility:** Given identical inputs, produces identical outputs
- **Objectivity:** No discretionary decisions affect boundaries
- **Auditability:** The complete partition hierarchy shows how each district formed
- **Speed:** All 50 states partition in hours, not months

Critics argue that redistricting inherently involves political tradeoffs that algorithms cannot navigate. We acknowledge this perspective but note two responses:

First, algorithmic transparency makes tradeoffs explicit. Rather than hiding decisions in back-room negotiations, the algorithm’s criteria are publicly documented. Communities can debate whether those criteria are appropriate, but the application remains consistent and verifiable.

Second, human judgment does not guarantee better outcomes. Our results show edge-weighted redistricting exceeds enacted district compactness in 37 of 50 states. In states with partisan-controlled redistricting, human judgment often produces less compact districts specifically because mapmakers optimize for partisan advantage rather than compactness.

7.2 From Baseline to Edge-Weighted: A Natural Progression

The 56% compactness improvement from baseline to edge-weighted mode demonstrates that algorithmic enhancements can bridge the performance gap with human mapmaking. Baseline recursive bisection provides a transparent, reproducible foundation that satisfies constitutional requirements. Edge-weighted enhancement adds compactness optimization without sacrificing transparency or reproducibility.

This progression suggests a path forward for redistricting reform:

1. **Baseline establishes viability:** Unweighted recursive bisection proves that algorithmic redistricting can satisfy constitutional requirements while eliminating partisan manipulation
2. **Edge-weighting adds competitiveness:** Boundary length optimization achieves compactness scores exceeding enacted districts
3. **Future enhancements remain possible:** Additional criteria (communities of interest, political subdivisions) could be incorporated while maintaining algorithmic transparency

The key insight is that algorithmic redistricting need not be static. Just as Huntington-Hill improved upon earlier apportionment methods, edge-weighted recursive bisection improves upon baseline recursive bisection. Future enhancements might further optimize for additional redistricting criteria while preserving the core advantages of transparency and objectivity.

7.3 When Algorithms Excel vs. Human Expertise

Our results reveal systematic patterns in where edge-weighted redistricting outperforms enacted districts:

Partisan-Drawn States: Illinois (+174.8%), Louisiana (+104.0%), New Hampshire (+101.7%), Texas (+85.7%), Massachusetts (+83.6%). In states where partisan legislatures control redistricting, algorithmic approaches substantially exceed enacted compactness. This confirms that human mapmakers in these contexts prioritize partisan advantage over compactness.

Commission-Drawn States: Iowa (-26.1%), Michigan (-5.4%), Washington (-5.5%). States with independent commissions often achieve compactness comparable to or exceeding algorithmic results. Human expertise, when insulated from partisan pressure, can produce high-quality districts.

This pattern suggests a nuanced role for algorithmic redistricting: as a *backstop against manipulation* rather than universal replacement for human judgment. In states with partisan control, algorithms offer clear improvements. In states with strong institutional protections, human expertise may perform comparably.

However, institutions change. A state with an independent commission today may switch to partisan control after the next election. Algorithmic redistricting provides *immunity to institutional decay*. The algorithm’s operation remains constant regardless of political composition of state government.

7.4 Perimeter Minimization as Compactness Proxy

Edge-weighted recursive bisection optimizes Polsby-Popper scores by directly minimizing perimeter. This raises the question: does optimizing one metric produce districts that score well on other measures?

The answer is largely yes. Districts with minimized perimeter also achieve strong Reock scores (0.396 mean, +14.1% over enacted). This suggests that perimeter minimization produces genuinely compact districts rather than merely gaming one specific metric.

Geographic features constrain compactness regardless of method. Coastal states (Florida, California) and states with irregular boundaries (Michigan, Maryland) naturally score lower on all compactness metrics. Edge-weighted mode does not overcome

these geographic constraints but performs comparably to human-drawn districts that face identical constraints.

7.5 Computational Scalability and Block-Level Data

Current implementation uses census tract-level data (84,414 tracts nationwide). Future work could explore block-level data (11+ million blocks), potentially improving resolution. However, computational complexity and METIS scalability would need careful evaluation.

The advantage of tract-level data is manageable graph size: even California (9,213 tracts) partitions in under 5 minutes. Block-level graphs would be 100-150× larger, potentially requiring hierarchical approaches or parallel METIS.

An alternative approach: use blocks for edge weight computation but tracts for partitioning. This would provide high-resolution boundary lengths while maintaining computational efficiency. The trade-off between resolution and computation time merits further investigation.

7.6 Multi-Objective Optimization and Additional Criteria

Edge-weighted recursive bisection optimizes compactness subject to population balance and contiguity. Traditional redistricting considers additional criteria:

- **Communities of interest:** Keeping neighborhoods, cities, or regions together
- **Political subdivisions:** Respecting county/municipal boundaries
- **Minority representation:** Ensuring compliance with Voting Rights Act
- **Competitiveness:** Creating districts with balanced partisan lean

Incorporating these criteria while maintaining algorithmic transparency poses challenges. Some can be integrated naturally:

Political subdivisions: Penalize edges that cross county boundaries by increasing their weights. This encourages METIS to prefer cuts within counties rather than between them.

Communities of interest: Allow manual pre-clustering of tracts that should remain together (e.g., a city), then partition the meta-graph. This preserves algorithmic transparency while respecting predefined communities.

Minority representation: Post-hoc analysis can identify whether algorithmically-generated districts provide minority communities with electoral influence. If not, constraints could be added to ensure specific regions remain united.

The key principle is maintaining *algorithmic transparency*. Any additional criteria must be encoded explicitly and applied consistently. Subjective human judgment at the boundary-drawing stage reintroduces opportunities for manipulation.

7.7 Legal and Constitutional Considerations

The Supreme Court has upheld equal population and contiguity as mandatory requirements (*Reynolds v. Sims*). Compactness, while widely valued, lacks consistent enforcement. Some state constitutions require compact districts but provide no quantitative definition, leading to disputes over interpretation.

Edge-weighted recursive bisection satisfies equal population and contiguity requirements while optimizing a well-defined compactness metric (Polsby-Popper). This provides courts with a concrete, reproducible standard. Rather than debating whether a district is "compact enough," litigation could focus on whether the algorithm's criteria appropriately balance competing objectives.

The Voting Rights Act requires consideration of minority representation. While our algorithm does not explicitly optimize for this, it produces majority-minority districts in areas with geographically concentrated minority populations. Courts could evaluate whether algorithmic outcomes satisfy VRA requirements, potentially requiring adjustments in specific cases.

7.8 Political Feasibility and Adoption Pathways

Algorithmic redistricting faces political resistance precisely because it eliminates partisan manipulation. Incumbent politicians benefit from district lines that protect their seats. Neither party has incentive to adopt algorithmic redistricting when they control map drawing.

However, several adoption pathways exist:

Citizen initiatives: States with ballot initiative processes allow voters to mandate redistricting criteria. California, Michigan, and Colorado adopted independent commissions through initiatives. Similar measures could mandate algorithmic redistricting.

Court-ordered remedies: When courts strike down gerrymandered maps, they often appoint special masters to draw replacement districts. Algorithmic redistricting provides courts with an objective, defensible alternative to human-drawn remedies.

Commission adoption: Existing independent commissions could use algorithmic redistricting as a starting point, making adjustments for communities of interest while maintaining transparency about what changed and why.

Voluntary adoption: States with bipartisan consensus against gerrymandering might adopt algorithmic redistricting as a commitment mechanism—neither party can manipulate future maps because the algorithm operates identically regardless of which party controls government.

The most promising path may be gradual adoption: states experiment with algorithmic redistricting, evaluate results, and refine criteria based on experience. Early adopters could demonstrate feasibility, encouraging broader acceptance.

7.9 Cross-Census Evidence for Algorithmic Neutrality

The cross-census validation provides compelling evidence for algorithmic objectivity. Across two census cycles with vastly different political contexts:

2010: Peak Gerrymandering Era

- Partisan control of most state legislatures
- Advanced GIS enabling precise optimization
- Pre-reform institutional environment
- Minimal judicial oversight of partisan gerrymanders
- **Result:** Enacted PP = 0.225 (historically low)

2020: Post-Reform Era

- Independent commissions in several major states
- Court challenges struck down aggressive gerrymanders

- Increased transparency requirements
- Public scrutiny via online mapping tools
- **Result:** Enacted PP = 0.305 (+35.7% improvement)

Despite these dramatic differences in political environment, the algorithm produced nearly identical results: 0.320 (2010) vs. 0.353 (2020). This **10.3% algorithmic variation** is dwarfed by the **42.4% gap** between algorithmic and enacted in 2010, demonstrating that:

1. **Geography, not politics, drives algorithmic output:** Population distribution and state boundaries—not partisan opportunities—determine district shape
2. **Algorithms provide stable baselines:** The consistency enables quantitative measurement of gerrymandering severity and reform effectiveness
3. **Reform is measurable:** The 50% reduction in the gap from 2010 to 2020 quantifies the impact of institutional reforms
4. **Further improvement is achievable:** Even post-reform 2020 districts fall 15.8% short of algorithmic baseline, indicating room for continued progress

This pattern strongly supports using algorithmic redistricting not necessarily as the *implementation method* but as an *objective standard* for evaluating enacted plans. Just as the Huntington-Hill method provides a clear formula for apportionment, edge-weighted recursive bisection provides a clear standard for compactness.

States could be required to justify deviations from the algorithmic baseline, explaining why political or community considerations necessitate less compact districts. This framework preserves human judgment for legitimate tradeoffs while providing courts and voters with quantitative tools to detect manipulation.

7.10 Limitations and Future Work

This study has several limitations:

Cross-census validation: We validated algorithm consistency using both 2010 and 2020 Census data. The remarkably stable algorithmic results

(0.320 vs. 0.353, only 10.3% variation) across different population distributions and political environments demonstrate robustness. This addresses concerns about temporal sensitivity and confirms that geographic factors, not political context, drive algorithm performance.

METIS parameters: We used standard METIS flags (`-contig`, `-minconn`, `-ufactor=1.005`). Alternative parameters might improve results. Systematic parameter tuning could optimize performance.

Recursive bisection structure: The binary tree of splits imposes hierarchical structure that may not match natural geographic divisions. Alternative partitioning schemes (e.g., k-way METIS, region-growing) merit investigation.

Edge weight calibration: We use boundary lengths directly. Alternative weighting schemes (e.g., logarithmic transformation, normalized weights) might improve compactness while balancing other criteria.

Limited community of interest modeling: Current implementation does not encode communities of interest beyond what emerges naturally from compactness. Developing methods to incorporate community preferences while maintaining transparency remains an open problem.

Future work should address these limitations while preserving the core advantages of algorithmic redistricting: transparency, reproducibility, and objectivity.

8 Conclusion

Congressional redistricting shapes American democracy but remains vulnerable to partisan manipulation. While the Huntington-Hill method resolved seat apportionment disputes through mathematical objectivity, boundary drawing has lacked comparable algorithmic frameworks. This paper demonstrates that edge-weighted recursive bisection—a graph partitioning approach that directly minimizes district perimeter—can match or exceed human mapmaking in compactness while providing transparency, reproducibility, and partisan neutrality.

8.1 Principal Findings

Our national-scale evaluation across all 435 congressional districts in 50 states yields four principal findings:

1. Edge-weighted recursive bisection substantially exceeds enacted district compactness

Mean Polsby-Popper score: 0.367 (algorithmic) vs. 0.305 (enacted)—a 20% improvement. In 37 of 50 states, the algorithm produces more compact districts than current enacted boundaries. Particularly strong performance in partisan-drawn states: Illinois +174.8%, Louisiana +104.0%, New Hampshire +101.7%.

2. Boundary length weighting dramatically improves baseline recursive bisection

Edge-weighted mode achieves 56% higher compactness than baseline unweighted mode (0.367 vs. 0.235). This improvement demonstrates that incorporating geometric information into graph partitioning yields substantial gains without sacrificing algorithmic transparency or reproducibility.

3. Algorithmic redistricting produces partisan outcomes reflecting political geography, not manipulation

Despite zero partisan input, edge-weighted districts show 56.5% Democratic-leaning vs. 34.0% Republican-leaning—a pattern consistent with urban concentration of Democratic voters and rural dispersion of Republican voters. The algorithm cannot engage in gerrymandering because it cannot distinguish between partisan voters; boundaries emerge from geographic optimization alone.

4. The algorithm satisfies constitutional requirements while optimizing compactness

All 435 districts maintain strict population balance (mean deviation 2.82%, max 4.98%) and guaranteed contiguity (enforced by METIS `-contig` flag). This demonstrates that constitutional constraints and compactness optimization are compatible—there is no inherent tradeoff between equal population and compact districts.

8.2 Implications for Redistricting Reform

These findings suggest several implications for redistricting reform:

Algorithmic redistricting is viable: The technology exists to automatically generate congressional districts that satisfy constitutional requirements while exceeding enacted district compactness in the majority of states. Technical feasibility is no longer a barrier to adoption.

Edge-weighting provides crucial enhancement: Baseline recursive bisection establishes that algorithmic approaches can work, but edge-weighted enhancement demonstrates they can excel. The 56% compactness improvement shows that algorithmic sophistication matters—simple unweighted graph partitioning is insufficient, but incorporating geometric information produces competitive results.

Partisan-drawn states benefit most: States where partisan legislatures control redistricting show the largest algorithmic improvements, confirming that human mapmakers in these contexts prioritize partisan advantage over compactness. Independent commissions produce more competitive results, but algorithmic approaches provide insurance against institutional decay.

Transparency enables accountability: Unlike commission-drawn or legislatively-drawn maps, algorithmic redistricting provides complete transparency. Every split decision is documented; the partition hierarchy is public; and citizens can verify that the algorithm operated as specified. This transparency shifts debate from suspicions of hidden agendas to explicit discussion of appropriate criteria.

8.3 Toward Mathematical Objectivity in Redistricting

The Huntington-Hill method demonstrates that mathematical algorithms can resolve intractable political problems. When seat apportionment became contentious in the early 20th century, Congress eventually adopted a formula that applies consistently, treats all states equally, and produces results that no single party controls. Eight decades later, Huntington-Hill operates without serious challenge because mathematical objectivity commands broader acceptance than any partisan compromise.

Edge-weighted recursive bisection offers a comparable framework for redistricting. Like Huntington-Hill, it:

- **Applies consistently:** The same algorithm operates in all states
- **Treats regions equally:** No geographic areas receive preferential treatment
- **Produces verifiable results:** Citizens can check that the algorithm operated correctly
- **Resists manipulation:** No single party can control outcomes

- **Optimizes explicit criteria:** Perimeter minimization directly maximizes Polsby-Popper compactness

The path from current practice to algorithmic redistricting will not be immediate. Incumbent politicians benefit from district lines that protect their seats. Neither party has incentive to voluntarily relinquish control. However, citizen initiatives, court-ordered remedies, and independent commission adoption provide alternative pathways. As more states experiment with algorithmic approaches, evidence of superior compactness and partisan neutrality may shift public opinion toward mathematical objectivity.

8.4 Future Directions

This work establishes edge-weighted recursive bisection as a viable redistricting approach but leaves several directions for future research:

Multi-objective optimization: Incorporating additional criteria (communities of interest, political subdivisions, minority representation) while maintaining algorithmic transparency

Higher resolution data: Exploring block-level data (11+ million units) vs. tract-level (84,414 units) to assess resolution vs. computation trade-offs

Alternative partitioning schemes: Comparing recursive bisection against k-way METIS, region-growing, and other graph partitioning approaches

Longitudinal analysis: Applying the algorithm to historical census data (2010, 2000) to evaluate consistency across redistricting cycles

Sensitivity analysis: Systematically varying METIS parameters and edge weight calibration to optimize performance

Legal evaluation: Working with legal scholars and civil rights organizations to ensure algorithmic approaches satisfy Voting Rights Act requirements

8.5 Closing Remarks

Congressional redistricting shapes representation in the U.S. House of Representatives, influencing policy outcomes and democratic legitimacy. For too long, this process has been vulnerable to partisan manipulation, generating public distrust and litigation. Algorithmic redistricting offers an alternative: transparent, reproducible, and objective district bound-

aries that optimize compactness while satisfying constitutional requirements.

Our results demonstrate that edge-weighted recursive bisection can exceed human-drawn district compactness in the majority of states while providing stronger guarantees against partisan manipulation than any human-drawn plan. Just as Huntington-Hill brought mathematical objectivity to congressional apportionment, edge-weighted recursive bisection can bring algorithmic transparency to redistricting.

The choice facing reformers is not between perfect algorithms and perfect human judgment—both have limitations. Rather, the choice is between *transparent mathematical criteria* that apply consistently and *opaque human judgment* that invites suspicion of manipulation. In an era of deep political polarization and declining trust in institutions, algorithmic redistricting offers a path toward restoring confidence in the fairness of electoral systems.

The technology exists. The mathematics is sound. The results are compelling. What remains is political will to adopt redistricting frameworks that prioritize transparency, objectivity, and geometric compactness over partisan advantage. This paper provides evidence that such frameworks can produce districts that exceed current practice while eliminating opportunities for gerrymandering. The question is not whether algorithmic redistricting can work—it demonstrably can. The question is whether we will choose mathematical objectivity over continued partisan manipulation.

Acknowledgments

We thank the U.S. Census Bureau for providing comprehensive redistricting data, the MIT Election Data + Science Lab for election results, and the Census Bureau’s TIGER/Line program for enacted district boundaries. This research was conducted using the METIS graph partitioning library developed by George Karypis and Vipin Kumar.



Figure 5: Histograms showing distribution of compactness improvements. Left: Edge-weighted vs. baseline (mean +56.0%, 43 of 50 positive). Right: Edge-weighted vs. enacted (mean +20.4%, 37 of 50 positive).



Figure 6: Cross-census comparison showing algorithmic consistency (blue bars, 0.32-0.35 range) vs. enacted variation (orange bars, 0.22-0.31 range). The gap between algorithmic and enacted narrows from 2010 to 2020, demonstrating measurable reform progress.