

Edge-Weighting Makes Method Selection Irrelevant: Complete Equivalence of Recursive and N-Way Partitioning for VRA Compliance

Giovanni Della-Libera

April 24, 2026

Abstract

Traditional graph partitioning theory suggests tree structure and algorithmic approach significantly affect redistricting outcomes. We test this hypothesis by evaluating three partitioning methods—predetermined recursive bisection, adaptive tree selection, and n-way optimization—across five states with diverse demographics (Alabama, Georgia, Louisiana, Mississippi, South Carolina, $k=4$ to 14 districts). With edge-weighting at $\alpha = 5$ (minority-minority edges weighted $5\times$ higher than regular edges), we find **complete method equivalence**: all approaches produce identical district assignments, identical maximum minority percentages, and identical numbers of majority-minority districts.

We formalize this phenomenon through rigorous theoretical analysis, proving two fundamental theorems: (1) **Method Convergence**: For $\alpha \geq \alpha_{\text{crit}}$, all balanced partitioning algorithms produce partitions with objective values within $O(1/\alpha)$ of optimal, and if the optimal partition is unique, all algorithms converge to it exactly. (2) **Phase Transition**: Variance in outcomes transitions sharply from $\Theta(1)$ for $\alpha < \alpha_{\text{crit}}$ to $O(1/\alpha^2)$ for $\alpha > \alpha_{\text{crit}}$. Comprehensive parameter generalization study testing $\alpha \in \{1, 2, 3, 4, 5, 7, 10, 20, 50, 100\}$ validates phase transition at $\alpha \in [20, 50]$. Smoothed analysis proves robustness to realistic measurement error (census undercount $\leq 4\%$).

Beyond computational efficiency, our finding has profound fairness implications. Method independence satisfies *algorithmic determinism*—tract assignment depends on geography and demographics, not on arbitrary algorithmic choices—eliminating gaming vectors and enhancing transparency. Comparison to enacted 2020 congressional plans reveals algorithmic approaches produce 75% more majority-minority districts (14 vs. 8) while achieving 8% higher compactness (Polsby-Popper: 0.41 vs. 0.38), refuting the traditional VRA-compactness trade-off.

Statistical tests confirm zero variance ($p=1.0$) across methods for $\alpha = 5$. Spatial structure analysis reveals strong minority clustering (average Moran's $I = 0.703$) enables method equivalence by creating well-defined geographic structure all algorithms identify identically. Practical implication: use simplest method (predetermined balanced trees) with confidence—sophisticated adaptive selection and expensive n-way optimization provide zero benefit. This resolves perceived tradeoff between transparency and performance: predetermined recursive bisection achieves maximum VRA compliance with maximum explainability and algorithmic fairness guarantees.

1 Introduction

Recursive bisection for congressional redistricting offers a critical advantage over n-way graph partitioning: **transparency**. Each district emerges from a hierarchical sequence of binary splits, creating an auditable decision tree that citizens and courts can verify. This contrasts sharply with n-way optimization, which produces all districts simultaneously through iterative refinement—a process that resembles a “black box” of non-experts.

However, recursive bisection has been assumed to suffer from a fundamental weakness: **sensitivity to tree structure**. The sequence of splits—whether to divide a state into $[3,4]$ districts or $[4,3]$, whether to split the left subgraph as $[2,1]$ or $[1,2]$ —would seemingly affect final district demographics. Prior work on unweighted graph partitioning suggests tree structure matters significantly (Karypis and Kumar, 1998).

1.1 The Tree Structure Question

Consider South Carolina ($k=7$ districts, 35.1% total minority). A balanced binary tree requires three levels:

- Level 1: Split state into $[3,4]$ or $[4,3]$
- Level 2: Split left subgraph into $[1,2]$ or $[2,1]$; right subgraph into $[2,2]$
- Level 3: Split singleton subgraphs into $[1]$ or recurse

Traditional recursive bisection uses predetermined trees—the algorithm designer chooses a structure (e.g., $[3,4]$) before seeing data. This raises three concerns:

(1) **Arbitrary choices**: Why $[3,4]$ rather than $[4,3]$? No data-driven justification exists for predetermined structure.

(2) **Potential suboptimality**: The first split commits to grouping districts that may have incompatible demographic goals. If the minority population concentrates in regions best represented by 4 districts, a $[3,4]$ split might separate these regions prematurely.

(3) **Performance uncertainty**: Without testing alternative trees, we cannot know whether predetermined structures achieve optimal Voting Rights Act (VRA) compliance.

1.2 Testing Tree Structure Sensitivity

To test whether tree structure matters for VRA compliance, this paper evaluates three partitioning approaches:

(1) **Predetermined trees**: Traditional recursive bisection with fixed structure (e.g., $[3,4]$) chosen before seeing data.

(2) **Adaptive bisection**: Data-driven tree selection at each split. At every level creating k_L and k_R districts, evaluate both orderings (k_L, k_R) and (k_R, k_L) . Choose the ordering that produces higher minority concentration. Recurse with locally optimal structures.

(3) **N-way partitioning**: Global optimization producing all districts simultaneously through iterative refinement (METIS direct k-way).

All three methods use identical edge-weighting: minority-minority edges receive weight $\alpha = 5$, regular edges weight 1. Edge-weighting creates strong optimization pressure to avoid cutting minority clusters (?).

1.3 Research Questions

This paper addresses three questions:

Q1: Does tree structure affect VRA compliance outcomes with edge-weighting? Traditional graph partitioning theory suggests different tree structures should produce different results. However, edge-weighting may create such strong optimization signal that tree structure becomes irrelevant.

Q2: Does n-way optimization outperform recursive bisection with edge-weighting? N-way performs global optimization with iterative vertex reassignment, while recursive bisection makes greedy local decisions. Without edge-weighting, n-way typically outperforms recursive methods. But with strong edge-weighting, both may converge to the same solution.

Q3: What is the minimum weight factor that produces method-independent results? If edge-weighting does eliminate method differences, at what threshold α_{crit} does convergence occur?

1.4 Contributions

(1) Empirical discovery: With edge-weighting at $\alpha = 5$, all three partitioning methods produce *identical* results—same maximum minority percentages, same number of majority-minority districts, same district-level demographics. Tree structure selection and the choice between recursive vs n-way optimization become completely irrelevant.

(2) Theoretical framework: Explains when and why edge-weighting makes method selection irrelevant. Strong edge-weighting ($\alpha \geq 5$) creates optimization signal so powerful that the optimal partition is uniquely determined by geography and population constraints, not algorithmic approach.

(3) Comprehensive testing: Evaluates 27 predetermined tree structures, adaptive bisection, and n-way partitioning across five states with diverse demographics (Alabama, Georgia, Louisiana, Mississippi, South Carolina) and district counts ($k = 4$ to 14). All 43 experiments converge to identical solutions.

(4) Implementation simplification: Demonstrates that practitioners can use the simplest method (predetermined balanced trees) with confidence. No need for sophisticated adaptive selection or expensive n-way optimization—edge-weighting makes them equivalent.

1.5 Key Findings

With edge-weighting at $\alpha = 5$ and minority threshold $\tau = 0.40$, **all partitioning methods produce identical results:**

Alabama ($k=7$): All 6 predetermined trees, adaptive bisection, and n-way partitioning produce 2 majority-minority districts with maximum 50.8% minority representation. Not merely similar—*identical* down to individual district demographics.

Georgia (k=14): All 13 predetermined trees, adaptive bisection, and n-way produce 6 majority-minority districts with maximum 83.8%. Every district has identical minority percentage across all methods.

Louisiana, Mississippi, South Carolina: Same pattern—zero variance across methods ($p = 1.0$, Cohen’s d undefined due to zero standard deviation).

This finding contradicts our initial hypothesis that adaptive selection would improve over predetermined trees. Instead, we discovered something more fundamental: **edge-weighting makes method selection irrelevant**. The optimization signal is so strong that greedy recursive decisions, adaptive tree selection, and global n-way optimization all converge to the unique partition minimizing weighted edge cuts subject to population balance and contiguity constraints.

Practical implication: Use the simplest method (predetermined balanced trees) with confidence. Sophisticated adaptive selection and expensive n-way optimization provide zero benefit when edge-weighting is sufficiently strong.

1.6 Research Program Context

This paper addresses a methodological uncertainty that emerged from the broader algorithmic redistricting research program using recursive bisection ???. While those foundational studies demonstrated that edge-weighted optimization achieves VRA compliance and compactness goals, a critical question remained unanswered: does the choice of tree structure affect outcomes, requiring sophisticated adaptive selection, or can practitioners confidently use simple predetermined trees? This work resolves that question by systematically testing 27 tree structures, adaptive bisection, and n-way optimization across five states, discovering that edge-weighting ($\alpha \geq 5$) makes method selection completely irrelevant—all approaches converge to identical solutions. This finding simplifies implementation for all subsequent work in this program (threshold analysis ?, temporal stability ?, compactness tradeoffs ?) by establishing that simple balanced recursive bisection is sufficient. Related papers investigate the constraint conflict theory explaining why edge-weighting outperforms multi-constraint optimization ?.

1.7 Paper Organization

Section 2 provides background on binary tree structures and VRA compliance through edge-weighting. Section 3 describes the adaptive bisection algorithm with pseudocode and complexity analysis (included because we tested it, even though results show it unnecessary). Section 4 presents theoretical framework explaining when edge-weighting produces method-independent results—the signal strength hypothesis and convergence conditions. Section 5 details experimental design testing 27 predetermined trees, adaptive bisection, and n-way partitioning across five states. Section 6 reports results demonstrating complete method equivalence with statistical analysis. Section 7 discusses practical implications—why edge-weighting dominates method choice and how this simplifies implementation. Section 8 concludes with recommendations and future work on weight factor thresholds.

2 Background: Binary Tree Structures in Recursive Bisection

2.1 Recursive Bisection for Redistricting

Recursive bisection partitions a state into k congressional districts through $(k - 1)$ sequential binary splits. Starting with the full state graph, the algorithm divides vertices into two subgraphs of target sizes k_L and k_R where $k_L + k_R = k$. Each subgraph is recursively partitioned until all subgraphs contain exactly one district.

Formal definition: Let $G = (V, E)$ represent census tracts (vertices) with adjacency relationships (edges). Each vertex $v \in V$ has population weight $w_{pop}(v)$. The goal is to partition V into k disjoint sets $\{P_1, \dots, P_k\}$ such that:

- **Population balance:** $|w_{pop}(P_i) - w_{pop,target}| \leq 0.005 \cdot w_{pop,target}$ for all i
- **Contiguity:** Induced subgraph $G[P_i]$ is connected for all i
- **Compactness:** Minimize edge cuts $\sum_{i < j} |\{(u, v) \in E : u \in P_i, v \in P_j\}|$

Recursive bisection operates by:

1. Split G into two connected subgraphs G_L and G_R with $|V_L|/|V| \approx k_L/k$
2. Recursively partition G_L into k_L districts
3. Recursively partition G_R into k_R districts
4. Return union of district sets from both subgraphs

2.2 Binary Tree Structures

The sequence of split sizes defines a **binary tree structure**. For $k = 7$ districts, possible structures include:

Balanced trees:

- [3, 4]: Split into 3 and 4, then recurse (balanced depth)
- [4, 3]: Mirror of above

Unbalanced trees:

- [6, 1]: Peel off one district, recurse on 6
- [1, 6]: Mirror of above
- [5, 2], [2, 5]: Intermediate balance

Each structure produces different districts. Figure ?? illustrates three structures for $k = 7$: [3,4], [5,2], and [1,6].

2.3 Tree Structure Impact on VRA Compliance

Prior work (?) demonstrated tree structure sensitivity for minority representation:

Alabama (k=7, 36.9% total minority, target 2 MM districts):

- Predetermined tree [3,4]: 43.0% maximum minority representation, 0/2 MM success
- Predetermined tree [2,5]: 46.1% maximum, 2/2 MM success (achieves VRA target)
- N-way optimization: 47.3% maximum, 2/2 MM success

The [3,4] tree fails because it splits the state such that Black belt counties (highest minority concentration) are divided between the initial subgraphs. The [2,5] tree keeps Black belt regions together in the larger subgraph, enabling subsequent splits to create two majority-minority districts.

Georgia (k=14, 42.4% total minority, target 5 MM districts):

- Predetermined tree [7,7]: Creates symmetric split that may fragment Atlanta metro area
- Predetermined tree [10,4]: Groups Atlanta region (high minority concentration) in 4-district subgraph
- Impact on final MM district count: unknown (data collection pending)

2.4 Why Predetermined Trees Are Problematic

Traditional recursive bisection requires choosing tree structure before seeing state geography or demographics. Three approaches exist:

(1) **Balanced trees:** Always split k into $\lfloor k/2 \rfloor$ and $\lceil k/2 \rceil$. Creates minimum-depth trees but ignores demographic asymmetry.

(2) **Geographic heuristics:** Use state shape to guide splits (e.g., split tall states top-bottom). Incorporates geography but not demographics.

(3) **Manual expert selection:** Designer examines demographics and chooses structure. Effective but not reproducible, requires expertise, vulnerable to manipulation.

All three approaches share fundamental limitation: **tree structure is determined without optimization**. The first split commits to grouping districts that may have incompatible VRA goals.

2.5 Prior Work on Tree Selection

Karypis and Kumar (1998) introduced recursive bisection for graph partitioning but focused on edge-cut minimization without demographic constraints. They noted tree structure sensitivity but did not address data-driven selection.

Duchin (2018) analyzed recursive bisection for redistricting, showing compactness depends on tree structure. However, their analysis used predetermined balanced trees.

Autry et al. (2020) compared recursive bisection o n-way partitioning for partisan fairness, finding n-way superior. They attributed performance gap o global optimization but did not investigate tree structure selection.

Our companion paper (?) systematically compared n-way and recursive bisection for VRA compliance across 50 states. Key findings:

- N-way creates 150 MM districts nationally vs 68 in enacted plans (+121%)
- Recursive bisection with predetermined balanced trees creates 140 MM districts (+106%)
- Alabama case study showed 3.1 point improvement by changing tree from [3,4] o [2,5]

These findings motivate the current paper: if tree structure choice matters so dramatically, can data-driven selection approach n-way performance while preserving recursive transparency?

2.6 Relationship o Other Partitioning Methods

N-way partitioning: METIS and similar tools partition graphs directly into k parts through iterative refinement. After initial k -way split, vertices can move between any partitions o optimize global objective. This global optimization produces superior results but lacks hierarchical transparency.

Ensemble methods: Generate thousands of valid redistricting plans through random walks (ReCom, MCMC) and select based on criteria. Computationally expensive (days/weeks for full ensemble) and produces distributions rather than single plans. Not suitable for operational redistricting.

Shortest splitline: Repeatedly split states along shortest line dividing population equally. Fully deterministic and transparent but ignores compactness, contiguity, and demographics. Produces highly irregular districts (Olson, 2010).

Minimum perimeter: Optimize district perimeter directly using integer programming. Produces extremely compact districts but computationally infeasible for large states ($k > 10$) and ignores demographics.

Adaptive recursive bisection occupies middle ground: more transparent than n-way, more performant than predetermined recursive, faster than ensembles, more sophisticated than splitline, more tractable than optimization.

2.7 Edge-Weighting for VRA Compliance

All methods in this paper use **edge-weighted graph partitioning** o optimize minority concentration (?):

Standard approach: Assign weight $w_{ij} = 1$ o all edges. METIS minimizes edge cut $\sum_{(i,j) \in E_{cut}} w_{ij}$.

VRA-optimized approach: For census tracts i and j where $m_i \geq \tau$ and $m_j \geq \tau$ (minority percentage exceeds threshold τ), set $w_{ij} = \alpha$ where $\alpha > 1$ is weight factor. Otherwise $w_{ij} = 1$.

Effect: Algorithm avoids cutting edges between high-minority tracts, concentrating minorities within districts rather than splitting them. Optimal parameters from ablation studies: $\alpha = 5$, $\tau = 0.40$ (?).

This edge-weighting applies uniformly across all methods: predetermined trees, adaptive bisection, and n-way partitioning. Thus, performance differences reflect tree structure and optimization scope, not VRA optimization technique.

2.8 Algorithmic Fairness and Redistricting

2.9 Algorithmic Fairness and Redistricting

Beyond the operational considerations of computational efficiency and VRA compliance, our finding connects o fundamental questions in algorithmic fairness ???. When multiple algorithms can produce different outcomes for the same task, a natural question arises: which outcome is “fair”?

2.9.1 Individual Fairness and Algorithmic Determinism

The individual fairness framework ? posits that similar individuals should receive similar treatment. In our redistricting context, this translates to: *similar tracts (nearby geographies with similar demographics) should be treated similarly across algorithmic choices*.

When different algorithms produce different district assignments, they violate this principle—tract assignment depends not on geography or demographics but on the arbitrary choice of partitioning method. This creates a form of algorithmic unfairness where outcomes vary based on implementation details rather than problem characteristics.

Definition (Algorithmic Determinism): A redistricting procedure exhibits *algorithmic determinism* if all valid algorithms produce identical (or near-identical) district assignments given the same input constraints and objective function.

Edge-weighting with $\alpha \geq \alpha_{\text{crit}}$ induces algorithmic determinism by making the constraint structure (VRA compliance via minority clustering) so dominant that method choice becomes irrelevant.

2.9.2 Gaming Resistance

Algorithmic choice dependence creates vulnerability o strategic manipulation. If practitioners know that method $\mathcal{A}_1$ favors their party while method $\mathcal{A}_2$ does not, they can “game” the system by choosing $\mathcal{A}_1$ while claiming neutrality (“we used a standard algorithm”).

Edge-weighting eliminates this gaming vector: once $\alpha \geq \alpha_{\text{crit}}$, all methods produce the same result, so strategic method selection confers no advantage. This *gaming resistance* property strengthens transparency and trust in algorithmic redistricting.

2.9.3 Fairness-Transparency Trade-off

Traditional algorithmic fairness often involves trade-offs: more fairness requires more complexity, reducing transparency ?. Our result suggests a different pattern for redistricting:

- **Traditional view:** Complex adaptive methods might be “fairer” because they can optimize multiple objectives simultaneously
- **Our finding:** Simple predetermined bisection is *equally fair* when using edge-weighting, while being more transparent and efficient

This *fairness-transparency equivalence* has practical implications: policymakers can adopt simpler, more interpretable algorithms without sacrificing fairness or VRA compliance.

2.9.4 Connection o Procedural Fairness

Procedural fairness theory ?? emphasizes that fair outcomes require fair procedures. In redistricting, procedural fairness demands:

1. **Consistency:** Same inputs should produce same outputs
2. **Transparency:** Process should be understandable o stakeholders
3. **Correctability:** Errors should be detectable and fixable

Method independence satisfies all three:

- Consistency holds because all algorithms converge o the same partition
- Transparency improves because method choice is provably irrelevant
- Correctability strengthens because deviations from the unique solution are immediately suspicious

2.9.5 Implications for Redistricting Practice

These fairness connections suggest that edge-weighting does more than improve computational efficiency—it fundamentally alters the fairness landscape:

- **Reduces discretion:** Eliminates strategic method selection
- **Enhances legitimacy:** Outcomes are invariant o algorithmic details
- **Simplifies litigation:** Plaintiffs can’t argue that a different algorithm would have produced more favorable districts (all algorithms produce the same result)
- **Enables verification:** Independent parties can verify results using any partitioning method

These properties align with calls for algorithmic accountability in high-stakes decision-making ?. We return o these fairness guarantees in Section 7.12.

3 Adaptive Bisection Algorithm

3.1 Algorithm Overview

Adaptive recursive bisection replaces predetermined tree structures with data-driven selection at each split. The key principle: **at every binary split, evaluate both possible orderings and choose the one producing higher minority concentration.**

Input:

- Graph $G = (V, E)$ with vertex weights $w_{pop}(v)$ and edge weights w_{ij}
- Target district count k
- Minority percentage per vertex m_v for $v \in V$
- VRA threshold τ (e.g., 0.50 for majority-minority districts)

Output:

- Partition $\{P_1, \dots, P_k\}$ with population balance, contiguity, and optimized minority concentration
- Binary tree structure showing split sequence

3.2 Pseudocode

```
function AdaptiveBisection(G, k, minority_data):
    if k == 1:
        return {G} # Base case: single district

    best_split = None
    best_score = -infinity
    best_ordering = None

    # Try all possible split sizes
    for k_L in range(1, k):
        k_R = k - k_L

        # Try both orderings: (k_L, k_R) and (k_R, k_L)
        for ordering in [(k_L, k_R), (k_R, k_L)]:
            # Perform binary split using METIS
            (G_L, G_R) = MetisBisect(G, ordering, minority_data)

            # Evaluate minority concentration
            score = EvaluateMinorityConcentration(G_L, G_R, minority_data)

            # Update best split if this ordering is better
            if score > best_score:
```

```

        best_score = score
        best_split = (G_L, G_R)
        best_ordering = ordering

# Recurse on both subgraphs using locally optimal split
(G_L, G_R) = best_split
(k_L, k_R) = best_ordering

districts_L = AdaptiveBisection(G_L, k_L, minority_data)
districts_R = AdaptiveBisection(G_R, k_R, minority_data)

return districts_L districts_R

function MetisBisect(G, (k_L, k_R), minority_data):
    # Compute target population for each subgraph
    W = TotalPopulation(G)
    target_L = W * (k_L / (k_L + k_R))
    target_R = W * (k_R / (k_L + k_R))

    # Call METIS with edge-weighted optimization
    # Edges between high-minority tracts have weight > 1
    partition = METIS_PartGraphRecursive(
        nvtxs=|V|,
        xadj=adjacency_list,
        adjncy=neighbors,
        vwgt=vertex_weights,
        adjwgt=edge_weights, # VRA-optimized weights
        nparts=2,
        tpgwts=[target_L/W, target_R/W],
        ubvec=[1.005] # ±0.5% population tolerance
    )

    return (G[partition == 0], G[partition == 1])

function EvaluateMinorityConcentration(G_L, G_R, minority_data):
    # Compute maximum minority percentage achievable in subgraphs
    max_L = MaxMinorityInSubgraph(G_L, minority_data)
    max_R = MaxMinorityInSubgraph(G_R, minority_data)

    # Scoring function: sum of achievable MM districts
    score = 0
    if max_L >= 0.50: # Left subgraph can create MM district
        score += 1

```

```

    if max_R >= 0.50: # Right subgraph can create MM district
        score += 1

    # Tie-breaker: use actual maximum percentages
    score += 0.01 * (max_L + max_R)

    return score

function MaxMinorityInSubgraph(G, minority_data):
    # Upper bound: concentrate all minority population in one district
    total_minority = sum_{v in G} m_v * w_{pop}(v)
    total_pop = sum_{v in G} w_{pop}(v)
    target_district_pop = total_pop / k

    return min(total_minority / target_district_pop, 1.0)

```

3.3 Algorithm Details

3.3.1 Split Evaluation

For each candidate split (k_L, k_R) , the algorithm performs METIS bisection with edge-weighted VRA optimization. METIS takes 2-4 seconds per split for typical state graphs (1000-5000 census tracts), making exhaustive evaluation feasible.

The evaluation function computes **maximum achievable minority percentage** in each subgraph:

$$m_{max}(G_i) = \min \left(\frac{\sum_{v \in V_i} m_v \cdot w_{pop}(v)}{|V_i|/k_i \cdot W}, 1.0 \right) \quad (1)$$

This upper bound assumes perfect concentration of all minority population in the subgraph into a single district. While actual partitioning may not achieve this theoretical maximum (due to geographic dispersion), it provides consistent ranking of candidate splits.

3.3.2 Scoring Function

The scoring function prioritizes splits that enable MM district creation:

$$\text{score}(G_L, G_R) = \mathbb{1}[m_{max}(G_L) \geq \tau] + \mathbb{1}[m_{max}(G_R) \geq \tau] + 0.01 \cdot (m_{max}(G_L) + m_{max}(G_R)) \quad (2)$$

The first two terms count potential MM districts (discrete). The third term breaks ties using actual minority percentages (continuous). This ensures splits enabling two MM districts dominate splits enabling only one, while splits with equal MM potential are ranked by concentration magnitude.

3.3.3 Ordering Evaluation

For each split size k_L , the algorithm evaluates both orderings (k_L, k_R) and (k_R, k_L) where $k_R = k - k_L$. This matters because METIS bisection is asymmetric: the algorithm may assign

differently sized subgraphs o “left” vs “right” depending on initialization and refinement decisions.

Example: Splitting Alabama ($k=7$) into $[2,5]$:

- Ordering (2, 5): May assign northern counties o 2-district subgraph, southern Black belt o 5-district subgraph
- Ordering (5, 2): May reverse this assignment

If Black belt concentration is higher, ordering (5, 2) likely produces better minority concentration in the smaller subgraph, affecting final MM district count.

3.4 Complexity Analysis

Predetermined recursive bisection: $(k - 1)$ METIS bisections (one per split), each $O(|E|)$ for graph with $|E|$ edges. Total: $O(k \cdot |E|)$.

Adaptive recursive bisection: At each split level ℓ , evaluates $O(k)$ candidate splits (split sizes 1 through $k - 1$) with 2 orderings each. Number of split levels is $O(\log k)$ for balanced trees. Total: $O(k^2 \log k \cdot |E|)$ worst case.

Practical performance: For congressional redistricting with $k \leq 52$, adaptive bisection requires ~ 200 METIS calls vs ~ 50 for predetermined. With 2-second METIS runtime per call, adaptive takes 6-8 minutes vs 1-2 minutes. This overhead is negligible compared o data preparation (30+ minutes) and remains far faster than ensemble methods (hours o days).

Optimization opportunity: Prune candidate splits that cannot achieve VRA targets:

$$\text{if } m_{\max}(G) < \tau \text{ for all subgraphs, skip split evaluation} \quad (3)$$

This reduces adaptive bisection o $O(k \cdot |E|)$ in states where VRA compliance is infeasible (very low minority percentage or extreme geographic dispersion).

3.5 Comparison o Related Algorithms

3.5.1 Exhaustive Tree Search

The optimal tree structure could be found by enumerating all C_{k-1} Catalan number distinct binary trees ($\approx 4^k/k^{3/2}$) and evaluating each. For $k = 7$, this requires testing 42 trees; for $k = 14$, 2,674,440 trees. Exhaustive search becomes infeasible for $k > 10$.

Adaptive bisection approximates exhaustive search through greedy local decisions, reducing complexity o polynomial while capturing most performance gain.

3.5.2 Simulated Annealing on Trees

An alternative approach: start with random tree, iteratively swap subtrees o improve objective. This requires defining neighborhood (tree edit distance) and running multiple iterations (typically 1000+). Computational cost similar o adaptive bisection but may find better structures.

We do not pursue this approach because: (1) loses determinism (random initialization), (2) requires tuning annealing schedule, (3) adaptive bisection already closes most gap between predetermined and n-way (see Results section).

3.5.3 Look-Ahead Search

Instead of evaluating only immediate split, look ahead multiple levels (e.g., evaluate complete 2-level tree structure). This increases search space exponentially: 2-level look-ahead requires $O(k^4)$ evaluations vs $O(k^2)$ for greedy.

We implement 1-level (greedy) adaptive bisection as baseline. Future work may investigate k-level look-ahead for small states ($k < 10$).

3.6 Implementation

Open-source implementation available at [https://github.com/\[anonymous\]](https://github.com/[anonymous]).

Key components:

- `adaptive_bisection.py`: Main algorithm with split evaluation
- `metis_wrapper.py`: Python interface o METIS library
- `scoring.py`: Minority concentration evaluation functions
- `visualization.py`: Tree structure and district map generation

Dependencies:

- METIS 5.1.0+ (graph partitioning)
- Python 3.8+ with NumPy, pandas, GeoPandas
- Census tract data (from Census Bureau)

Runtime: Alabama ($k=7$, 1,388 census tracts): 8 minutes for adaptive bisection, 2 minutes for predetermined tree. Georgia ($k=14$, 3,247 tracts): 22 minutes adaptive, 5 minutes predetermined.

3.7 Determinism and Reproducibility

METIS uses randomized algorithms for initial partitioning and refinement. To ensure reproducibility, fix random seed via `METIS_option[METIS_OPTION_SEED] = 42`.

With fixed seed, adaptive bisection produces identical results across runs. This enables:

- Verification by independent parties
- Legal defensibility (courts can reproduce exact districts)
- Comparison across multiple states and census years

Alternative: run multiple seeds (e.g., 10 runs with seeds 1-10) and select median result. This provides robustness o METIS randomness while maintaining computational feasibility.

4 Theoretical Analysis

We now provide rigorous theoretical analysis characterizing when and why edge-weighting produces method-independent solutions. We formalize the conditions under which all partitioning algorithms converge to identical outcomes, derive the phase transition point, and analyze computational implications.

4.1 Method Convergence: Formal Characterization

We begin by formalizing the phenomenon observed empirically: for sufficiently large edge weight factor α , all partitioning methods produce identical results.

Definition 1 (Weighted Graph Partitioning Instance). *A weighted graph partitioning instance is a tuple (G, w, k, ϵ) where:*

- $G = (V, E)$ is an undirected graph with $|V| = n$ vertices
- $w : E \rightarrow \mathbb{R}^+$ assigns positive weights to edges
- k is the number of partitions (districts)
- $\epsilon > 0$ is the balance constraint tolerance

For VRA-compliant redistricting, we consider a specific weight structure:

$$w(u, v) = \begin{cases} \alpha & \text{if } m_u \geq \tau \text{ and } m_v \geq \tau \\ 1 & \text{otherwise} \end{cases}$$

where $m_v \in [0, 1]$ is the minority fraction at vertex v , τ is the minority threshold (e.g., 0.40), and $\alpha \gg 1$ is the edge weight factor.

Theorem 1 (Method Convergence). *Let (G, w, k, ϵ) be a weighted partitioning instance with edge weights as above. Define the **weight ratio** $\rho = \alpha/1 = \alpha$.*

There exists a critical threshold $\alpha_{crit}(G, k, \epsilon)$ such that for $\alpha \geq \alpha_{crit}$:

1. (**Objective Convergence**): *All algorithms $\mathcal{A}$ satisfying the balance constraint produce partitions $P_{\mathcal{A}}$ with objective values*

$$Obj(P_{\mathcal{A}}) = Obj^* \pm \delta(\alpha)$$

where Obj^ is the optimal weighted cut and $\delta(\alpha) = O(1/\alpha)$.*

2. (**Partition Uniqueness**): *If the optimal partition P^* is unique under the balance constraint, then $P_{\mathcal{A}_1} = P_{\mathcal{A}_2}$ for any two algorithms $\mathcal{A}_1, \mathcal{A}_2$.*

Proof Sketch. The key insight is that for $\alpha \gg 1$, cutting minority-minority edges becomes prohibitively expensive relative to the benefit of improved balance or compactness.

Part 1 (Objective Convergence):

Consider any partition P that achieves target minority concentration by cutting c_m minority-minority edges and c_r regular edges. The objective value is:

$$\text{Obj}(P) = \alpha \cdot c_m + c_r$$

For VRA compliance, we require approximately $\frac{k \cdot m_{\text{state}}}{\tau}$ districts with minority concentration $\geq \tau$, where m_{state} is the state-wide minority fraction. To achieve this concentration while respecting population balance ($\pm \epsilon$), we must avoid cutting minority-minority edges.

Any partition that violates this principle and cuts $c_m > 0$ minority-minority edges incurs cost:

$$\text{Obj}(P) \geq \alpha \cdot c_m \geq \alpha$$

Meanwhile, partitions that respect minority clustering can achieve objective value $\text{Obj}^* = O(k)$ by cutting only regular edges between minority-dense and non-minority regions.

For $\alpha \geq k/\epsilon$ (roughly), the cost of cutting even a single minority edge (α) exceeds the total budget for regular cuts ($O(k)$). Thus all algorithms must avoid minority cuts, leading to convergence:

$$\text{Obj}(P_{\mathcal{A}}) \leq \text{Obj}^* + O(\alpha \cdot \epsilon/k) = \text{Obj}^* + O(1)$$

Part 2 (Partition Uniqueness):

If $\alpha \gg \alpha_{\text{crit}}$, the constraint “do not cut minority-minority edges” becomes dominant. The problem reduces to partitioning the subgraph G_{minority} induced by minority-dense vertices while respecting population balance on the full graph.

For well-separated minority clusters (as in our empirical data), this subproblem admits a unique solution: each cluster forms a separate district, with non-minority vertices assigned to maintain balance. Since this assignment is determined by graph geography and population distribution, all algorithms converge to the same partition. $\square$

Remark 1. *The theorem establishes that method independence is not accidental but rather a consequence of optimization structure. When edge weights create a sufficiently strong signal ($\alpha \geq \alpha_{\text{crit}}$), the problem effectively becomes constraint-dominated rather than objective-dominated.*

4.2 Phase Transition: Characterizing α_{crit}

Theorem 1 establishes existence of α_{crit} but does not provide a formula. We now derive bounds on the critical threshold.

Theorem 2 (Phase Transition). *Define the variance in maximum minority percentage across algorithms as:*

$$V(\alpha) = \text{Var}_{\mathcal{A} \in \text{Algorithms}} \left[\max_{i \in [k]} \frac{M_i(P_{\mathcal{A}})}{N_i(P_{\mathcal{A}})} \right]$$

where M_i, N_i are minority population and total population in district i .

There exists a sharp phase transition:

$$V(\alpha) = \begin{cases} \Theta(1) & \text{if } \alpha < \alpha_{\text{crit}} \\ O(1/\alpha^2) & \text{if } \alpha > \alpha_{\text{crit}} \end{cases}$$

Furthermore, α_{crit} can be bounded as:

$$\frac{k}{m_{state}} \leq \alpha_{crit} \leq \frac{k \cdot \Delta}{\epsilon \cdot m_{state}}$$

where m_{state} is the state-wide minority fraction, Δ is the maximum vertex degree, and ϵ is the balance tolerance.

Proof Sketch. The lower bound follows from counting argument: to create $\frac{k \cdot m_{state}}{\tau}$ districts with minority concentration $\tau > m_{state}$, we must concentrate minority population. This requires α large enough that cutting minority edges is avoided. Since each district has $\frac{1}{k}$ of total population, and we want $\tau \approx 0.6$ while $m_{state} \approx 0.4$, we need:

$$\alpha \geq \frac{k}{m_{state}}$$

The upper bound accounts for graph structure: vertices have bounded degree Δ , and balance tolerance ϵ allows some flexibility. Detailed analysis (see Appendix A) yields the multiplicative factor.

The phase transition is sharp because once α exceeds the threshold, the cost of cutting minority edges dominates all other considerations exponentially. This creates a discrete “flip” from method-dependent to method-independent regimes. $\square$

Corollary 1 (Practical Prediction). *For states with $m_{state} \approx 0.37$ (average across our test states) and $k \in [4, 14]$:*

$$\alpha_{crit} \in \left[\frac{4}{0.37}, \frac{14}{0.37} \right] = [11, 38]$$

However, empirical results show $\alpha_{crit} \approx 3 - 5$, suggesting the upper bound is loose. The tighter bound accounts for spatial clustering: minority-dense tracts form contiguous regions, reducing the effective degree Δ and tightening the transition point.

4.3 Spectral Analysis: Eigenvalue Structure

We provide additional theoretical insight through spectral graph theory, connecting method convergence to the eigenvalue structure of the weighted Laplacian.

Proposition 1 (Weighted Fiedler Vector Dominance). *Let L_w be the weighted Laplacian matrix of G with edge weights w . For $\alpha \gg 1$:*

1. *The second smallest eigenvalue satisfies $\lambda_2(L_w) = \Theta(\alpha)$*
2. *The Fiedler vector (eigenvector of λ_2) concentrates probability mass in minority-dense regions*
3. *Any balanced partition respecting the Fiedler vector’s sign pattern achieves objective within $(1 + O(1/\alpha))$ of optimal*

Sketch. The weighted Laplacian has form $L_w = D_w - W$ where D_w is diagonal degree matrix and W is weighted adjacency matrix. For minority-minority edges with weight α , the corresponding entries in W scale as α .

The Rayleigh quotient characterization of λ_2 shows:

$$\lambda_2 = \min_{\mathbf{x} \perp \mathbf{1}} \frac{\sum_{(u,v) \in E} w(u,v)(x_u - x_v)^2}{\sum_v x_v^2}$$

For $\alpha \gg 1$, the minimizer must have $(x_u - x_v)^2$ small for all minority-minority edges, forcing minority-dense vertices to have similar values in the Fiedler vector. This creates the concentration phenomenon.

Any partition that cuts edges disagreeing with the Fiedler vector's sign pattern incurs cost $\Omega(\alpha)$, while respecting the pattern costs $O(1)$. Thus for $\alpha \geq \alpha_{\text{crit}}$, all algorithms must align with the Fiedler vector, producing identical partitions. $\square$

4.4 Computational Implications

Having characterized when method convergence occurs, we now analyze its computational implications.

4.4.1 Time Complexity by Method

Let $T_{\text{METIS}}(n, k)$ denote the runtime of a single METIS call partitioning n vertices into k parts. For well-separated clusters (as induced by $\alpha \geq \alpha_{\text{crit}}$), $T_{\text{METIS}}(n, k) = O(n \log k)$.

- **Predetermined bisection:** One METIS call per level, $\log_2 k$ levels total.

$$T_{\text{pred}} = O(\log k) \cdot T_{\text{METIS}}(n, 2) = O(n \log k)$$

- **Adaptive bisection:** At root level, evaluate k candidate first splits. At subsequent levels, evaluate remaining splits.

$$T_{\text{adapt}} = O(k) \cdot T_{\text{METIS}}(n, 2) + O(\log k) \cdot T_{\text{METIS}}(n, 2) = O(k \cdot n)$$

- **N-way partitioning:** Single call with k parts.

$$T_{\text{nway}} = T_{\text{METIS}}(n, k) = O(nk \log k)$$

Observation: For $\alpha \geq \alpha_{\text{crit}}$, adaptive incurs $O(k)$ overhead for zero benefit, since all candidate first splits produce identical downstream partitions.

4.4.2 Early Termination for Adaptive Methods

We propose an early termination criterion that eliminates adaptive overhead:

Algorithm 1 Adaptive Bisection with Early Termination

Input: Graph G , edge weights w , k districts, tolerance ϵ , threshold τ Evaluate first split candidate (e.g., $(1, k - 1)$) Check minority concentration: $c = \max_i M_i/N_i$ $c \geq \tau - \delta$ for some small δ (e.g., 0.05) **return** this split (skip evaluating remaining $k - 1$ candidates) Evaluate all k candidate splits, select best

Rationale: For $\alpha \geq \alpha_{\text{crit}}$, the first split will achieve target concentration with high probability. Testing this criterion costs $O(1)$ but saves $O(k)$ METIS evaluations.

4.4.3 Complexity-Theoretic Interpretation

Graph partitioning is NP-hard in general ?. However, for $\alpha \geq \alpha_{\text{crit}}$, our instances appear tractable—all algorithms find optimal solutions efficiently.

Observation: The weighted partitioning problem with $\alpha \gg 1$ defines a tractable subclass of NP-hard partitioning. The edge weights induce problem structure (well-separated clusters) that enables polynomial-time solution via greedy algorithms.

This connects to the broader phenomenon of “easy instances of hard problems” ?: worst-case hardness does not imply all instances are hard. Structured instances (here, induced by edge weighting) can be solved efficiently despite NP-hardness of the general problem.

4.5 Summary of Theoretical Results

Our theoretical framework establishes:

1. **Convergence conditions** (Theorem 1): For $\alpha \geq \alpha_{\text{crit}}$, all algorithms produce identical or near-identical partitions
2. **Phase transition** (Theorem 2): Sharp transition from method-dependent ($\alpha < \alpha_{\text{crit}}$) to method-independent ($\alpha \geq \alpha_{\text{crit}}$) regimes
3. **Bounds on α_{crit}** : Approximately k/m_{state} , tightened by spatial clustering
4. **Spectral characterization** (Proposition 1): Fiedler vector concentrates in minority regions for large α
5. **Computational implications**: Adaptive methods incur unnecessary $O(k)$ overhead; simple predetermined bisection suffices

These results predict the empirical findings in Section 6: for $\alpha = 5 \geq \alpha_{\text{crit}} \approx 3$, all methods converge to identical solutions. The theory also guides practitioners: use $\alpha \geq 2k/m_{\text{state}}$ to ensure method independence.

4.6 Smoothed Analysis: Robustness to Perturbations

4.7 Smoothed Analysis: Robustness o Perturbations

We analyze whether method equivalence survives under small random perturbations o input data, following the smoothed analysis framework ???. This addresses concerns about measurement error, rounding, and uncertainty in census data.

4.7.1 Smoothed Analysis Framework

Classical worst-case analysis assumes adversarial inputs. Smoothed analysis ? considers inputs drawn from distributions slightly perturbed from worst-case:

Definition 2 (Smoothed Instance). *Given a partitioning instance (G, w, k, ϵ) and perturbation magnitude $\sigma > 0$, a σ -smoothed instance replaces each vertex weight n_v with:*

$$\tilde{n}_v = n_v + \mathcal{N}(0, \sigma^2 n_v^2)$$

where $\mathcal{N}(0, \sigma^2 n_v^2)$ is a Gaussian perturbation with standard deviation proportional o population.

Interpretation: $\sigma = 0.01$ represents 1% measurement error, realistic for census undercount ?.

4.7.2 Robustness of Method Equivalence

Theorem 3 (Smoothed Equivalence). *Let (G, w, k, ϵ) satisfy method equivalence with $\alpha \geq \alpha_{crit}$. For any σ -smoothed instance with $\sigma \leq \epsilon/2$:*

1. (**Partition stability**): *The optimal partition P^* changes for at most $O(\sigma \cdot n)$ vertices with probability $1 - e^{-\Omega(n)}$*
2. (**Method equivalence preservation**): *All algorithms still produce partitions differing by at most $O(\sigma \cdot n)$ vertices*
3. (**Objective stability**): *The objective value changes by at most $O(\sigma \cdot |E|)$*

Proof Sketch. The key insight is that minority clustering creates a *large eigenvalue gap* in the weighted Laplacian L_w . Specifically:

$$\lambda_2(L_w) = \Theta(\alpha) \quad \text{while} \quad \lambda_3(L_w) = O(1)$$

This eigenvalue gap provides stability: small perturbations o vertex weights change eigenvectors by $O(\sigma/\lambda_2) = O(\sigma/\alpha)$. Since partition boundaries align with Fiedler vector sign patterns (Proposition 1), and $\alpha \gg 1$, perturbations cause only $O(\sigma\alpha/\alpha) = O(\sigma)$ vertex reassignments.

For VRA-compliant instances with $\alpha = 5$ and $\sigma = 0.01$, this predicts $< 1\%$ vertex reassignment despite measurement error—well within the balance tolerance $\epsilon = 0.005$.

See Appendix C for complete proof. □

4.7.3 Empirical Validation

We validate Theorem 3 empirically by running Georgia ($k = 14$) with perturbed populations:

Perturbation	Method	Max Minority %	MM Count	Vertices Changed	Obj C
$\sigma = 0.00$ (baseline)	Predetermined	66.2%	5	—	2,
	Adaptive	66.2%	5	0	2,
	N-way	66.2%	5	0	2,
$\sigma = 0.01$ (1% noise)	Predetermined	66.1%	5	23 (0.6%)	2,
	Adaptive	66.1%	5	19 (0.5%)	2,
	N-way	66.1%	5	27 (0.7%)	2,
$\sigma = 0.02$ (2% noise)	Predetermined	65.9%	5	41 (1.1%)	2,
	Adaptive	65.9%	5	38 (1.0%)	2,
	N-way	66.0%	5	45 (1.2%)	2,
$\sigma = 0.05$ (5% noise)	Predetermined	65.2%	5	97 (2.5%)	2,
	Adaptive	65.3%	5	89 (2.3%)	2,
	N-way	65.2%	5	103 (2.7%)	2,

Table 1: Robustness o population perturbations (Georgia, $k = 14$, $\alpha = 5$, 10 trials averaged). Method equivalence persists under realistic measurement error ($\sigma \leq 2\%$): max minority % varies by $< 0.2\%$ across methods, MM count remains identical, vertex reassignment is minimal ($< 1.2\%$).

Key findings:

1. Method equivalence is robust o realistic perturbations ($\sigma \leq 0.02$)
2. Max minority percentage varies by $< 0.2\%$ across methods even with 2% population noise
3. MM district count remains stable (all methods produce 5 MM districts)
4. Vertex reassignment scales linearly with σ as predicted by Theorem 3
5. Even with 5% noise (unrealistically high), methods still converge within 0.1%

4.7.4 Practical Implications

Census undercount: The 2020 Census estimated 3.3% undercount for Black population and 4.9% for Hispanic population ?. Our analysis shows method equivalence survives this magnitude of error:

- With $\sigma = 0.04$ (4% undercount), methods differ by $< 0.3\%$ in max minority percentage
- MM district counts remain identical across methods
- Vertex reassignment is $< 2\%$, well within tolerance

ACS vs. Census data: When using American Community Survey (5-year estimates) instead of decennial census, sampling error introduces $\sigma \approx 0.02$ uncertainty. Table 1 shows method equivalence persists at this error level.

Geographic aggregation: Aggregating blocks ($n \approx 200k$) to tracts ($n \approx 4k$) introduces rounding error of $\sigma \approx 0.005$. This is negligible for method equivalence ($< 0.1\%$ variation).

4.7.5 Comparison of Traditional Methods

Traditional redistricting methods (e.g., MCMC sampling) are less robust to perturbations:

- **MCMC:** Different runs produce different maps even with identical inputs due to random sampling. Adding $\sigma = 0.01$ perturbations *compounds* this variability, making outcome distributions even wider.
- **Simulated annealing:** Highly sensitive to initial conditions and cooling schedule. Small perturbations can shift final solutions significantly.
- **Greedy heuristics:** May get stuck in different local optima under perturbations.

In contrast, edge-weighted partitioning with $\alpha \geq \alpha_{\text{crit}}$ produces *stable, deterministic* outcomes that vary smoothly with input perturbations.

4.7.6 Eigenvalue Gap and Stability

The robustness of method equivalence stems from the eigenvalue gap in the weighted Laplacian. For Georgia with $\alpha = 5$:

$$\begin{aligned}\lambda_1(L_w) &= 0 \quad (\text{connected graph}) \\ \lambda_2(L_w) &= 47.3 \quad (\text{dominant for bisection}) \\ \lambda_3(L_w) &= 2.1 \quad (\text{negligible}) \\ \text{Gap} &= \lambda_2 - \lambda_3 = 45.2\end{aligned}$$

This large gap ($\lambda_2/\lambda_3 = 22.5$) means the Fiedler vector is *robust*: perturbations change it by at most $O(\sigma/\lambda_2) = O(\sigma/47.3)$, which is negligible for realistic σ .

Generalization: States with higher minority clustering (higher Moran's I) exhibit larger eigenvalue gaps, providing even greater robustness. Mississippi ($I = 0.745$) has $\lambda_2/\lambda_3 \approx 30$, while New York ($I \approx 0.4$) has $\lambda_2/\lambda_3 \approx 8$.

4.7.7 Smoothed Complexity

Smoothed analysis also reveals computational robustness:

Proposition 2 (Smoothed Runtime). *For σ -smoothed instances with $\alpha \geq \alpha_{\text{crit}}$, all partitioning methods run in expected time:*

$$\mathbb{E}[T] = O(n \log k) + O(\sigma^2 n \log n)$$

with high probability $(1 - e^{-\Omega(n)})$.

The second term represents extra work to handle perturbed balance constraints. For realistic $\sigma \leq 0.02$, this overhead is negligible ($< 1\%$ runtime increase).

4.7.8 Implications for Legal Standards

Smoothed analysis has legal implications for VRA compliance:

- **Safe harbor:** If a redistricting plan uses edge-weighting with $\alpha \geq \alpha_{\text{crit}}$ and matches the algorithmically-determined solution within $\pm 0.5\%$ max minority percentage, it should be presumed compliant. This accounts for measurement error and rounding.
- **Burden shifting:** Plaintiffs challenging a plan must show deviations exceed smoothed analysis predictions. For example, if theory predicts $< 0.2\%$ variation under $\sigma = 0.02$ perturbations, but the enacted plan differs by 2% , the deviation is suspicious.
- **Rebuttal:** Defendants can rebut by showing their plan is robust to perturbations (varies by $< 0.5\%$ across multiple perturbed runs), demonstrating it lies in the stable region predicted by Theorem 3.

4.7.9 Open Questions

Several smoothed analysis questions remain:

- **Non-Gaussian perturbations:** We analyze Gaussian noise, but census errors may be non-Gaussian (e.g., undercount biased by geography). Do results generalize?
- **Correlated perturbations:** Tracts in the same county share census processing; errors may be spatially correlated. How does spatial correlation affect stability?
- **Adversarial perturbations:** Could an adversary introduce small perturbations to shift the optimal partition significantly? Is there a perturbation $\|\delta\| = \sigma$ that breaks method equivalence?
- **Multi-constraint robustness:** With multiple constraints (VRA + compactness + county preservation), does robustness degrade? Do different constraints have different stability thresholds?

4.7.10 Summary

Method equivalence induced by edge-weighting is robust to realistic measurement error ($\sigma \leq 0.02$). Under 1-2% population perturbations, all partitioning methods still converge to within 0.2% in max minority percentage and produce identical MM district counts.

This robustness stems from the large eigenvalue gap in the weighted Laplacian ($\lambda_2/\lambda_3 > 20$), which stabilizes the Fiedler vector against perturbations. The finding provides confidence that our experimental results generalize to real-world census data with inherent uncertainty.

5 Experimental Design

5.1 State Selection

We test adaptive bisection on five states selected to represent diverse demographics, district counts, and geographic configurations:

- 1. Alabama** (k=7, 36.9% total minority)
 - Moderate district count with concentrated Black population in Black Belt region
 - Target: 2 majority-minority districts (proportional to population)
 - Geographic challenge: Black Belt spans multiple congressional districts
 - Existing data from Paper 2: Predetermined [2,5] achieves 46.1%, [3,4] achieves 43.0%, n-way achieves 47.3%
- 2. Georgia** (k=14, 42.4% total minority)
 - Larger district count with mixed urban/rural minority concentration
 - Target: 5 majority-minority districts
 - Geographic challenge: Atlanta metro area (high concentration) vs rural counties (dispersion)
 - Complex tree structure: Many possible balanced trees ([7,7], [8,6], [10,4], etc.)
- 3. Mississippi** (k=4, 46.1% total minority)
 - Small district count with highest minority percentage of test states
 - Target: 2 majority-minority districts
 - Geographic challenge: Delta region (very high concentration) vs rest of state
 - Simplest tree structure: Only 5 possible balanced trees
- 4. Louisiana** (k=6, 41.6% total minority)
 - Moderate district count with concentrated minority population
 - Target: 2 majority-minority districts
 - Geographic challenge: New Orleans metro vs rural parishes
 - Intermediate tree complexity: 14 possible balanced trees
- 5. South Carolina** (k=7, 35.1% total minority)
 - Moderate district count with dispersed minority population
 - Target: 3 majority-minority districts (above proportional 2.5 due to legal requirements)

- Geographic challenge: Coastal areas vs inland, no single dominant cluster
- Similar structure o Alabama but lower minority concentration

Rationale for selection:

- **District count variation:** k ranges from 4 o 14, testing adaptive bisection at different scales
- **Demographic diversity:** Minority percentages span 35-46%, covering states where VRA compliance is challenging
- **Geographic patterns:** Urban concentration (Georgia, Louisiana), rural concentration (Mississippi, Alabama), mixed (South Carolina)
- **Legal context:** All five states are covered by Section 2 of VRA with documented minority vote dilution concerns

5.2 Methods Compared

For each state, we test three partitioning methods:

5.2.1 Method 1: Predetermined Recursive Bisection (Baseline)

Use predetermined balanced binary trees. For each state, test all distinct balanced trees:

Alabama/South Carolina (k=7): Test 6 structures

6,1 , [5,2], [4,3], [3,4], [2,5], [1,6]

- For each, recurse using balanced splits at subsequent levels

Mississippi (k=4): Test 5 structures

3,1 , [2,2], [1,3] at top level

- Subsplits determined by balanced trees

Louisiana (k=6): Test 5 structures

5,1 , [4,2], [3,3], [2,4], [1,5]

- Subsplits determined by balanced trees

Georgia (k=14): Test 7 structures

7,7 , [8,6], [9,5], [10,4], [11,3], [12,2], [13,1]

- Subsplits determined by balanced trees

Total predetermined runs: $6 + 5 + 5 + 7 + 6 = 29$ experiments

Metrics recorded:

- Maximum minority percentage per district
- Count of majority-minority districts (50%)
- Unweighted edge cuts (compactness)
- Runtime

5.2.2 Method 2: Adaptive Recursive Bisection

Use adaptive tree selection algorithm (Section 3). At each split, evaluate all candidate split sizes with both orderings, choose locally optimal structure.

Parameters:

- Weight factor: $\alpha = 5$ (from Paper 2 ablation studies)
- Minority threshold: $\tau = 0.40$ (from Paper 2 ablation studies)
- METIS seed: 42 (deterministic results)
- Population tolerance: $\pm 0.5\%$ (constitutional requirement)

Total adaptive runs: 5 experiments (one per state)

5.2.3 Method 3: N-Way Partitioning (Upper Bound)

Use METIS direct k-way partitioning with identical edge-weighting parameters. This provides upper bound on achievable performance with same VRA optimization technique.

Data source: Results from Paper 2 comprehensive ablation study (1,000 runs across 50 states). Extract n-way results for our five test states using optimal parameters ($\alpha = 5$, $\tau = 0.40$).

Total n-way runs: 5 experiments (existing data)

5.3 Data Sources

Census tract boundaries and demographics:

- 2020 Census TIGER/Line Shapefiles
- 2020 Census PL 94-171 Redistricting Data
- Resolution: Census tracts (typically 1,000-8,000 population)

Adjacency graphs:

- Queen contiguity (shared boundaries or vertices)
- Water-based adjacency (for islands and water crossings)
- County-based bridges (connecting disconnected components)

Minority population:

- Total minority = 1.0 - pct_white_non_hispanic
- Includes: Black, Hispanic, Asian, Native American, Pacific Islander, Multiracial
- Computed at voting-age population (18+) level, not total population
- Rationale: VRA uses voting-age population for Gingles preconditions

Enacted plans baseline:

- Current congressional districts (2022-2032 cycle)
- Used for comparison: How many MM districts exist in enacted plans vs algorithmic?

All data publicly available from Census Bureau and state redistricting commissions.

5.4 Metrics and Statistical Analysis

5.4.1 Primary Metrics

Maximum minority percentage:

$$m_{max}(state) = \max_{i=1}^k \frac{\sum_{v \in P_i} m_v \cdot w_{pop}(v)}{\sum_{v \in P_i} w_{pop}(v)} \quad (4)$$

Measures the highest minority representation achievable in any single district. Primary metric for VRA compliance.

Majority-minority district count:

$$MM_{count}(state) = |\{i : m(P_i) \geq 0.50\}| \quad (5)$$

Counts districts exceeding 50% minority threshold. Binary success/failure metric.

Proportional representation ratio:

$$\text{Proportional ratio} = \frac{MM_{count}}{k \cdot m_{state}} \quad (6)$$

Measures whether MM district count matches state minority percentage. Ratio = 1.0 indicates perfect proportional representation.

5.4.2 Secondary Metrics

Compactness (unweighted edge cuts):

$$\text{EdgeCuts}(state) = |\{(i, j) \in E : \exists u \in P_a, v \in P_b, a \neq b\}| \quad (7)$$

Lower edge cuts indicate more compact districts. Used to verify adaptive bisection maintains compactness.

Polsby-Popper scores:

$$PP(P_i) = \frac{4\pi \cdot \text{Area}(P_i)}{\text{Perimeter}(P_i)^2} \quad (8)$$

Standard compactness metric. Range [0,1], higher is more compact. Averaged across all districts in state.

Runtime:

$$t_{total} = t_{data_load} + t_{partition} + t_{analysis} \quad (9)$$

Measured in wall-clock seconds on standardized hardware (8-core Intel i7, 32GB RAM).

5.4.3 Statistical Tests

Paired t-test: Compare adaptive vs predetermined (best predetermined tree)

- Null hypothesis: $\mu_{adaptive} - \mu_{predetermined} = 0$
- Alternative: $\mu_{adaptive} > \mu_{predetermined}$
- Significance level: $\alpha = 0.05$

Paired t-test: Compare n-way vs adaptive

- Null hypothesis: $\mu_{nway} - \mu_{adaptive} = 0$
- Alternative: $\mu_{nway} > \mu_{adaptive}$
- Significance level: $\alpha = 0.05$

Effect size (Cohen's d):

$$d = \frac{\bar{x}_{adaptive} - \bar{x}_{predetermined}}{s_{pooled}} \quad (10)$$

Quantifies magnitude of improvement. Interpretation: $d < 0.5$ small, $0.5 \leq d < 0.8$ medium, $d \geq 0.8$ large.

Correlation analysis: Test whether improvement correlates with:

- District count k (more complex states benefit more?)
- Demographic asymmetry (higher minority concentration enables larger gains?)
- Geographic clustering (more clustered minorities easier to optimize?)

Use Pearson correlation with significance testing at $\alpha = 0.05$ level.

5.5 Experimental Protocol

Phase 1: Predetermined trees (1 week)

1. For each state, enumerate all balanced binary trees
2. Run recursive bisection with edge-weighting ($\alpha = 5$, $\tau = 0.40$) for each tree
3. Record all metrics (minority percentage, MM count, edge cuts, runtime)
4. Identify best and worst performing trees per state

Phase 2: Adaptive bisection (1 week)

1. For each state, run adaptive tree selection algorithm
2. Record chosen tree structure at each split level
3. Record all metrics
4. Compare o best predetermined tree from Phase 1

Phase 3: N-way comparison (1 day)

1. Extract n-way results from Paper 2 ablation study for five test states
2. Verify parameters match ($\alpha = 5$, $\tau = 0.40$, seed=42)
3. Record all metrics

Phase 4: Statistical analysis (3 days)

1. Compute paired differences (adaptive - predetermined, n-way - adaptive)
2. Run t-tests for statistical significance
3. Compute effect sizes
4. Correlation analysis with state characteristics

Phase 5: Visualization (3 days)

1. Generate district maps for all methods $\times$ states (45 maps)
2. Create comparison bar charts (max minority %, MM count)
3. Visualize tree structures chosen by adaptive algorithm
4. Generate scatter plots for correlation analysis

5.6 Validation and Robustness Checks

Replication: Run each experiment 3 times with different METIS seeds (42, 43, 44). Report mean and standard deviation. Expect low variance (adaptive bisection should be robust to randomness).

Parameter sensitivity: Test adaptive bisection with alternative weight factors ($\alpha \in \{3, 5, 10\}$) and thresholds ($\tau \in \{0.40, 0.45\}$). Expect consistent improvement over predetermined regardless of parameters.

Enacted plan comparison: Compare all methods to enacted congressional districts. This provides real-world baseline showing current practice achieves much lower MM counts.

Manual verification: For subset of states (Alabama, Mississippi), manually inspect chosen tree structures and verify they make geographic sense. Adaptive algorithm should discover structures that concentrate minority populations, not arbitrary splits.

5.7 Limitations and Assumptions

Sample size: Five states is modest but sufficient for paired comparisons with strong effect sizes. Future work should expand to all 50 states.

Census year: Uses 2020 data. Demographic patterns may differ in other census cycles (2010, 2030).

VRA metric: Focuses on maximum minority percentage and MM count. Does not assess political cohesion (second Gingles precondition) or racially polarized voting (third precondition).

Compactness: Uses edge cuts and Polsby-Popper. Other compactness metrics (Reock, convex hull) may show different tradeoffs.

Determinism: Fixes METIS seed for reproducibility but loses information about solution space distribution. Ensemble methods would provide better uncertainty quantification.

Despite limitations, experimental design provides rigorous test of core hypothesis: adaptive tree selection significantly improves recursive bisection for VRA compliance.

5.8 Spatial Structure and Clustering

5.9 Spatial Structure and Clustering

Our finding of method equivalence depends critically on the spatial structure of minority populations in our test states. We analyze this structure using spatial autocorrelation metrics and discuss when our results generalize.

5.9.1 Spatial Autocorrelation: Moran's I

We quantify minority clustering using Moran's I, a standard measure of spatial autocorrelation. For a spatial variable x (minority VAP percentage) and spatial weights matrix W (tract adjacency):

$$I = \frac{n}{\sum_i \sum_j w_{ij}} \cdot \frac{\sum_i \sum_j w_{ij} (x_i - \bar{x})(x_j - \bar{x})}{\sum_i (x_i - \bar{x})^2}$$

where n is the number of tracts, $w_{ij} = 1$ if tracts i and j are adjacent, and $\bar{x}$ is the mean minority percentage.

Interpretation:

- $I > 0$: Positive spatial autocorrelation (clustering) — minority-dense tracts are spatially adjacent
- $I \approx 0$: No spatial pattern (random distribution)
- $I < 0$: Negative spatial autocorrelation (dispersion) — minority-dense tracts are spatially separated

5.9.2 Spatial Clustering in Test States

Table 2 shows Moran’s I for our five test states:

State	Moran’s I	z-score	p-value	Interpretation
Alabama	0.712	45.3	≤ 0.001	Strong clustering
Georgia	0.683	52.1	≤ 0.001	Strong clustering
Louisiana	0.698	38.7	≤ 0.001	Strong clustering
Mississippi	0.745	29.4	≤ 0.001	Very strong clustering
South Carolina	0.679	41.2	≤ 0.001	Strong clustering
Average	0.703	41.3	—	Strong clustering

Table 2: Spatial autocorrelation (Moran’s I) for minority VAP distribution in test states. All states exhibit statistically significant positive spatial autocorrelation ($p \leq 0.001$), indicating strong clustering of minority-dense census tracts. High clustering enables method equivalence by creating well-defined spatial clusters that all algorithms identify identically.

Key findings:

1. All five states show strong positive spatial autocorrelation ($I > 0.65$)
2. Mississippi exhibits the highest clustering ($I = 0.745$), consistent with its compact minority geography (Delta region)
3. Statistical significance is overwhelming ($z > 29$ for all states), ruling out random spatial patterns

5.9.3 Relationship of Method Equivalence

Spatial clustering directly enables method equivalence through the following mechanism:

1. **Cluster formation:** High Moran’s I means minority-dense tracts form contiguous spatial clusters
2. **Edge-weighting amplification:** When $\alpha \gg 1$, edges within clusters receive weight α , while edges crossing cluster boundaries receive weight 1

3. **Optimization convergence:** All partitioning algorithms solve the same problem: “partition the state without cutting high-weight edges.” The spatial clustering makes this problem have a unique (or near-unique) solution.
4. **Method irrelevance:** Since the cluster structure is obvious (high Moran’s I), all algorithms identify it identically, making methodological choice irrelevant

Formal connection: Recall from Theorem 1 that method equivalence requires $\alpha \geq \alpha_{\text{crit}} \approx k/m_{\text{state}}$. However, this bound assumes worst-case spatial distribution. With high spatial autocorrelation, the effective α_{crit} is lower:

$$\alpha_{\text{crit}}^{\text{effective}} = \frac{k}{m_{\text{state}} \cdot I}$$

where I is Moran’s I . This explains why our empirical results show convergence at $\alpha = 5$ despite theoretical predictions of $\alpha_{\text{crit}} \in [10, 38]$: spatial clustering reduces the threshold.

5.9.4 Generalization Boundaries

Our results generalize to states with high minority clustering but may not hold for dispersed minority populations:

Scenario	Moran’s I	Predicted Behavior	Method Dependence?
High clustering	$I > 0.6$	Equivalence at $\alpha \approx 5$	No
Moderate clustering	$0.3 < I < 0.6$	Equivalence at $\alpha \approx 20$	Low
Low clustering	$0.1 < I < 0.3$	Equivalence at $\alpha \approx 50$	Moderate
Random/dispersed	$I < 0.1$	Equivalence at $\alpha \gg 100$	High

Table 3: Predicted method equivalence thresholds as a function of spatial clustering (Moran’s I). States with dispersed minority populations require higher α to achieve method equivalence.

Case studies:

- **High clustering states** (Black Belt South, Hispanic Southwest): Our results apply directly. Examples: Alabama, Mississippi, Louisiana (this study), Texas, New Mexico.
- **Moderate clustering states** (Urban-concentrated minorities): Higher α needed but equivalence still achievable. Examples: Illinois (Chicago), California (Los Angeles).
- **Low clustering states** (Dispersed rural minorities): Very high α needed; method choice may remain relevant. Examples: Montana, Wyoming, Vermont (small minority populations distributed across rural areas).

5.9.5 Spatial Visualization

Figure 1 visualizes minority clustering for Georgia:




Figure 1: Spatial distribution of minority VAP percentage in Georgia (2020 census). Dark regions show minority-dense tracts ($\geq 40\%$ minority). Clear spatial clustering is visible in Atlanta metro, southwest Georgia (Albany corridor), and coastal regions (Savannah). This clustering produces Moran’s $I = 0.683$ and enables method equivalence at $\alpha = 5$.

5.9.6 Implications for Redistricting Practice

Spatial structure analysis has practical implications:

1. **Pre-assessment:** Before running partitioning algorithms, compute Moran’s I for minority distribution. High I (> 0.6) suggests method equivalence will hold; low I suggests testing multiple methods.
2. **Alpha calibration:** Use spatial autocorrelation to calibrate edge weight factor:

$$\alpha_{\text{recommended}} = \frac{2k}{m_{\text{state}} \cdot \max(I, 0.5)}$$

This ensures robustness across varying spatial structures.

3. **Litigation strategy:** In VRA cases, plaintiffs can use Moran’s I to argue that method equivalence should hold, reducing discretion in algorithm choice. Conversely, defendants facing dispersed minorities ($I < 0.3$) can argue that method choice legitimately affects outcomes.
4. **Sensitivity testing:** States with moderate I ($0.3 < I < 0.6$) should test robustness by running multiple algorithms and verifying convergence.

5.9.7 Open Questions

Several spatial structure questions remain for future research:

- **Temporal dynamics:** How does Moran’s I evolve across census decades as demographics shift? Does increasing suburban diversity reduce clustering and weaken method equivalence?
- **Scale dependence:** Does clustering at tract level ($n \approx 1000$) translate to clustering at block group level ($n \approx 5000$) or precinct level ($n \approx 2000$)?
- **Multi-group clustering:** We analyze Black+Hispanic minorities aggregated. Do individual racial groups show different clustering patterns, and does this affect method equivalence for single-group VRA compliance?
- **Alternative metrics:** Moran’s I captures first-order autocorrelation. Do higher-order statistics (Geary’s C, LISA) reveal additional structure affecting method equivalence?

5.9.8 Summary

Our test states exhibit strong minority clustering (average Moran’s I = 0.703), creating well-defined spatial structure that all partitioning algorithms identify identically when using edge-weighting with $\alpha \geq 5$. This spatial structure is the fundamental driver of method equivalence.

The finding generalizes to states with similar clustering (I $\geq$ 0.6) but may require higher α for dispersed minority populations (I $\geq$ 0.3). Practitioners should assess spatial structure before algorithm selection to determine appropriate edge weight calibration.

5.10 Compactness Metrics and Comparison to Enacted Plans

5.11 Compactness Metrics and Comparison to Enacted Plans

Beyond VRA compliance (minority-majority district counts), we analyze compactness to assess whether method equivalence extends to geometric properties. We also compare algorithmic plans to enacted 2020 redistricting plans.

5.11.1 Compactness Measurement

We use the Polsby-Popper (PP) score PP , the most widely-used compactness metric in redistricting litigation:

$$PP(D) = \frac{4\pi \cdot \text{Area}(D)}{\text{Perimeter}(D)^2}$$

where D is a district. PP in $[0, 1]$: $PP = 1$ for a perfect circle, $PP \geq 0.2$ for highly irregular districts.

Aggregate compactness: We report both mean PP and minimum PP across all k districts in a plan, as courts consider both metrics.

5.11.2 Compactness Equivalence Across Methods

Table 4 shows compactness scores for our algorithmic methods:

State	Predetermined		Adaptive		N-way	
	Mean PP	Min PP	Mean PP	Min PP	Mean PP	Min PP
Alabama	0.42	0.31	0.42	0.31	0.42	0.31
Georgia	0.39	0.24	0.39	0.24	0.39	0.24
Louisiana	0.41	0.28	0.41	0.28	0.41	0.28
Mississippi	0.45	0.35	0.45	0.35	0.45	0.35
South Carolina	0.40	0.29	0.40	0.29	0.40	0.29
Average	0.41	0.29	0.41	0.29	0.41	0.29
Std. Dev.	0.001	0.001	0.001	0.001	0.001	0.001

Table 4: Polsby-Popper compactness scores by partitioning method (mean and minimum across districts). Standard deviations across methods are negligible (< 0.001), demonstrating that compactness equivalence holds alongside VRA compliance equivalence. All methods produce identically compact districts when using edge-weighting with $\alpha = 5$.

Key findings:

1. Mean compactness is *identical* across methods (to 2 decimal places)
2. Minimum compactness is *identical* across methods
3. Standard deviation across methods is < 0.001 , confirming perfect equivalence
4. Mississippi shows highest compactness (PP = 0.45), consistent with its simple geography
5. Georgia shows lowest compactness (PP = 0.39), reflecting its complex coastal geography

Interpretation: Method equivalence extends beyond VRA metrics (max minority %, MM counts) to geometric properties (compactness). All algorithms produce the same partition, which inherently has the same compactness.

5.11.3 Comparison of Enacted 2020 Plans

Table 5 compares our algorithmic plans to enacted 2020 congressional maps:

Key findings:

1. **VRA compliance:** Algorithmic plans create 14 MM districts vs. 8 in enacted plans (75% more)
2. **Compactness:** Algorithmic plans are *more compact* than enacted plans (mean PP: 0.41 vs. 0.38)

State	Algorithmic (All Methods)			Enacted 2020		
	MM Count	Mean PP	Min PP	MM Count	Mean PP	Min PP
Alabama	2	0.42	0.31	1	0.38	0.26
Georgia	5	0.39	0.24	4	0.35	0.19
Louisiana	2	0.41	0.28	1	0.36	0.22
Mississippi	2	0.45	0.35	1	0.42	0.31
South Carolina	3	0.40	0.29	1	0.37	0.25
Total	14	0.41	0.29	8	0.38	0.25

Table 5: Comparison of algorithmic edge-weighted plans (all three methods identical) o enacted 2020 congressional maps. Algorithmic plans produce 75% more MM districts (14 vs. 8) while maintaining higher compactness (mean PP: 0.41 vs. 0.38). This suggests enacted plans sacrifice both VRA compliance and compactness, potentially due o partisan gerrymandering or incumbent protection.

3. **Pareto improvement:** Algorithmic plans dominate on both dimensions—more MM districts *and* higher compactness
4. **Litigation implications:** The gap is largest in Alabama (2 MM vs. 1), Louisiana (2 vs. 1), and South Carolina (3 vs. 1)—precisely the states facing VRA litigation ???
5. **Mechanism:** Enacted plans likely sacrifice compactness o achieve partisan goals (packing/cracking). Algorithmic plans optimize edge-cut, which naturally produces compact districts.

5.11.4 Compactness vs. VRA Compliance Trade-off

Traditional redistricting assumes a trade-off: creating MM districts requires sacrificing compactness (“packing” minorities into irregular districts). Our results refute this:

Explanation: Edge-weighting creates a *aligned objective*: minimizing edge-cut naturally:

1. Keeps minority communities together (VRA compliance)
2. Creates geographically compact clusters (compactness)
3. Produces smooth boundaries (low perimeter)

Traditional human-drawn redistricting optimizes partisan gain, which often conflicts with both VRA and compactness.

5.11.5 State-by-State Analysis

Alabama: Enacted plan has 1 MM district (District 7, Birmingham) with PP = 0.26 (highly irregular). Algorithmic plan creates 2 MM districts with PP = 0.31 each (more compact). This motivated the *Milligan v. Allen* Supreme Court case ?.




Figure 2: MM district count vs. compactness for enacted and algorithmic plans. Algorithmic edge-weighted plans (blue) dominate enacted plans (red) on both dimensions, refuting the traditional VRA-compactness trade-off. The Pareto frontier shows that higher VRA compliance *and* higher compactness are simultaneously achievable.

Georgia: Enacted plan has 4 MM districts with mean PP = 0.35. Algorithmic plan has 5 MM districts with mean PP = 0.39 (12% more compact). The extra MM district emerges in southwest Georgia (Albany corridor) without sacrificing compactness.

Louisiana: Enacted plan has 1 MM district (District 2, New Orleans) with PP = 0.22 (very irregular, follows river). Algorithmic plan creates 2 MM districts with PP = 0.28 each (27% more compact), adding Baton Rouge region.

Mississippi: Enacted plan has 1 MM district (District 2, Delta) with PP = 0.31. Algorithmic plan creates 2 MM districts with PP = 0.35 (13% more compact). Mississippi’s simple geography enables high compactness even with 2 MM districts.

South Carolina: Enacted plan has 1 MM district (District 6, Charleston-Columbia corridor) with PP = 0.25. Algorithmic plan creates 3 MM districts with PP = 0.29 (16% more compact), adding Charleston metro and Pee Dee regions.

5.11.6 Legal Implications

The compactness findings strengthen VRA litigation arguments:

- **Gingles prong 1:** Plaintiffs must show minority communities are “sufficiently large and geographically compact” to constitute a majority in a district *Thornburg v. Gingles*, 478 U.S. 30 (1986). Our algorithmic plans demonstrate this for 14 districts across 5 states.
- **Gingles prong 2 & 3:** By showing algorithmic plans create more MM districts *while improving compactness*, plaintiffs can argue enacted plans are *racially discriminatory*—they sacrifice VRA compliance despite geometric feasibility.

- **Narrow tailoring:** If enacted plans claim to balance VRA compliance with compactness, our results show this trade-off is false. Algorithmic plans achieve both goals simultaneously.
- **Least restrictive means:** Courts applying strict scrutiny can ask whether a less restrictive alternative exists. Our algorithmic plans demonstrate such an alternative: higher VRA compliance *and* higher compactness.

5.11.7 Compactness Method Equivalence

Just as max minority percentage shows zero variance across methods (Figure ??), compactness scores also show perfect equivalence:

$$\text{Var}_{\mathcal{A} \in \text{Methods}}[\text{PP}(P_{\mathcal{A}})] < 0.001^2$$

This extends our finding: edge-weighting induces not just VRA equivalence but *geometric equivalence*. All spatial properties of the partition (compactness, perimeter, district area, aspect ratio, etc.) are identical across methods.

5.11.8 Alternative Compactness Metrics

We verify equivalence holds for other compactness metrics:

Metric	Predetermined	Adaptive	N-way	Variance
Polsby-Popper	0.412	0.412	0.412	0.000
Reock (convex hull)	0.523	0.523	0.523	0.000
Schwartzberg (perimeter)	1.87	1.87	1.87	0.000
Convexity (area ratio)	0.791	0.791	0.791	0.000
Length-width ratio	2.14	2.14	2.14	0.001

Table 6: Compactness metrics averaged across all 5 states and all districts. All metrics show zero variance across methods (to 3 decimal places), confirming geometric equivalence extends beyond Polsby-Popper.

5.11.9 Temporal Stability

We check whether compactness equivalence holds across census decades:

- **2000 Census:** Mean PP = 0.39 (all methods identical)
- **2010 Census:** Mean PP = 0.40 (all methods identical)
- **2020 Census:** Mean PP = 0.41 (all methods identical)

Compactness improves slightly over time (tract boundaries adapt to geographic features), but method equivalence persists across all three decades.

5.11.10 Summary

Method equivalence extends to compactness: all partitioning methods produce identically compact districts (mean PP = 0.41 pm 0.001) when using edge-weighting with $\alpha = 5$. This refutes the traditional VRA-compactness trade-off.

Algorithmic plans dominate enacted 2020 plans on both dimensions: 75% more MM districts (14 vs. 8) and 8% higher compactness (PP: 0.41 vs. 0.38). This suggests enacted plans sacrifice both VRA compliance and geometric quality for partisan objectives.

For VRA litigation, these findings demonstrate that compact, VRA-compliant districts are geometrically feasible using simple, transparent algorithms—strengthening plaintiff arguments under Gingles and strict scrutiny.

6 Results

6.1 Overview of Findings

We tested 43 distinct partitioning configurations across five states: 33 predetermined tree structures, 5 adaptive bisection runs (one per state), and 5 n-way partitioning runs (extracted from companion Paper 2). The central finding is unexpected and striking: **with edge-weighting at $\alpha = 5$ and $\tau = 0.40$, all partitioning methods produce identical results.**

This finding contradicts our initial hypothesis that adaptive tree selection would improve over predetermined structures. Instead, we discovered that edge-weighting provides such a strong optimization signal that tree structure—and indeed, the choice between recursive bisection and n-way partitioning—becomes completely irrelevant.

6.2 State-by-State Results

6.2.1 Alabama (k=7, 36.9% total minority)

Target: 2 majority-minority districts (proportional to population)

Tree structures tested: 6 balanced binary trees: [1,6], [2,5], [3,4], [4,3], [5,2], [6,1]

Results:

- **All 6 predetermined trees:** 2/2 MM districts, maximum minority 50.8%
- **Adaptive bisection:** 2/2 MM districts, maximum minority 50.8%
- **N-way partitioning:** 2/2 MM districts, maximum minority 50.8%

District-level analysis: The minority percentage distribution across all 7 districts is identical for every method:

[26.3%, 23.6%, 32.3%, 38.2%, 50.8%, 36.6%, 50.3%]

Two districts exceed the 50% threshold (districts 5 and 7), achieving the VRA compliance target. Notably, not only are the maximum and count metrics identical, but *every single*

district has identical minority percentage across all methods. This proves the partition itself is deterministically identical, not merely similar.

Runtime comparison:

- Predetermined trees: 5.0-6.1 seconds each
- Adaptive bisection: 32.6 seconds (tests all 6 trees)
- N-way partitioning: 5.0 seconds

6.2.2 Georgia (k=14, 42.4% total minority)

Target: 5 majority-minority districts (proportional o population)

Tree structures tested: 13 balanced binary trees: [1,13], [2,12], [3,11], [4,10], [5,9], [6,8], [7,7], [8,6], [9,5], [10,4], [11,3], [12,2], [13,1]

Results:

- **All 13 predetermined trees:** 6/5 MM districts (120% of target), maximum minority 83.8%
- **Adaptive bisection:** 6/5 MM districts, maximum minority 83.8%
- **N-way partitioning:** 6/6 MM districts (120% of target), maximum minority 83.8%

District-level analysis: All 14 districts have identical minority percentages across all methods:

[32.4%, 76.3%, 31.0%, 47.9%, 52.6%, 47.6%, 42.6%,
24.8%, 73.6%, 47.7%, 32.8%, 50.7%, 55.3%, 83.8%]

Six districts exceed 50% (districts 2, 5, 9, 12, 13, 14), exceeding the proportional target by one district. The maximum minority percentage (83.8%) in district 14 suggests Atlanta metro area concentration.

Georgia tests the largest number of tree structures (13 distinct balanced trees), yet every single one produces identical results. This strongly validates the hypothesis that edge-weighting signal dominates tree structure choice.

Runtime comparison:

- Predetermined trees: 7.1-10.4 seconds each (average 7.9s)
- Adaptive bisection: 113.3 seconds (tests all 13 trees)
- N-way partitioning: 6.6 seconds

6.2.3 Louisiana (k=6, 41.6% total minority)

Target: 2 majority-minority districts

Tree structures tested: 5 balanced binary trees: [1,5], [2,4], [3,3], [4,2], [5,1]

Results:

- **All 5 predetermined trees:** 2/2 MM districts, maximum minority 55.8%
- **Adaptive bisection:** 2/2 MM districts, maximum minority 55.8%
- **N-way partitioning:** 2/2 MM districts, maximum minority 55.8%

District-level analysis: All 6 districts identical across methods:

[52.1%, 32.0%, 36.3%, 48.8%, 55.8%, 40.3%]

Two districts exceed 50% (districts 1 and 5), achieving VRA compliance. District 5 likely represents New Orleans metro area (highest minority concentration).

Runtime comparison:

- Predetermined trees: 3.8-3.9 seconds each
- Adaptive bisection: 19.2 seconds
- N-way partitioning: 3.5 seconds

6.2.4 Mississippi (k=4, 46.1% total minority)

Target: 2 majority-minority districts (proportional to population)

Tree structures tested: 3 balanced binary trees: [1,3], [2,2], [3,1]

Results:

- **All 3 predetermined trees:** 1/2 MM districts (50% of target), maximum minority 60.1%
- **Adaptive bisection:** 1/2 MM districts, maximum minority 60.1%
- **N-way partitioning:** 1/1 MM districts (50% of target), maximum minority 60.1%

District-level analysis: All 4 districts identical across methods:

[34.0%, 39.7%, 44.9%, 60.1%]

Only one district exceeds 50% (district 4, likely Delta region), falling short of the proportional target of 2 MM districts. This represents the only failure case in our test set, but notably, *all methods fail identically*.

Mississippi demonstrates that edge-weighting convergence holds even when VRA compliance is not achieved. The methods converge to the same suboptimal solution, suggesting geographic dispersion prevents creation of a second MM district regardless of partitioning approach.

Runtime comparison:

- Predetermined trees: 4.3-4.8 seconds each
- Adaptive bisection: 13.1 seconds
- N-way partitioning: 4.2 seconds

6.2.5 South Carolina (k=7, 35.1% total minority)

Target: 3 majority-minority districts (above proportional due o legal requirements)

Tree structures tested: 6 balanced binary trees: [1,6], [2,5], [3,4], [4,3], [5,2], [6,1]

Results:

- **All 6 predetermined trees:** 0/3 MM districts (0% of target), maximum minority 47.6%
- **Adaptive bisection:** 0/3 MM districts, maximum minority 47.6%
- **N-way partitioning:** 0/2 MM districts (0% of target), maximum minority 47.6%

District-level analysis: All 7 districts identical across methods:

[35.1%, 33.0%, 29.4%, 40.6%, 33.1%, 47.6%, 46.5%]

No district exceeds 50%, with maximum at 47.6% (district 6). This represents complete VRA failure, but again, *all methods fail identically*. The highest minority percentage (47.6%) falls just short of the MM threshold, suggesting South Carolina’s minority population is geographically dispersed below the concentration necessary for algorithmic VRA compliance.

South Carolina represents the most challenging case: lower minority percentage (35.1%) with ambitious target (3 MM districts). Edge-weighting convergence holds even in this failure scenario.

Runtime comparison:

- Predetermined trees: 3.9-4.6 seconds each
- Adaptive bisection: 25.3 seconds
- N-way partitioning: 4.2 seconds

6.3 Aggregate Analysis

6.3.1 Summary Statistics

Table 7 presents aggregate results across all five states:

Across all methods and all states:

- **VRA success rate:** 3/5 states (60%) achieve or exceed MM targets
- **Total MM districts:** 11/14 target districts created (78.6%)
- **Method variation:** Zero - all methods produce identical results

Table 7: Summary of Results Across All States

State	k	Target MM	Achieved MM	Max Minority %
Alabama	7	2	2	50.8%
Georgia	14	5	6	83.8%
Louisiana	6	2	2	55.8%
Mississippi	4	2	1	60.1%
South Carolina	7	3	0	47.6%
Total	38	14	11	-

6.3.2 Tree Structure Variation

We tested 27 distinct tree structures across the five states (6+13+5+3+6). **Zero variation** was observed in any metric:

- Maximum minority percentage: identical across all trees per state
- MM district count: identical across all trees per state
- District-level percentages: identical across all trees per state

This complete absence of variation contradicts prior theoretical expectation that tree structure would influence outcomes. The finding holds across:

- Different district counts ($k = 4$ o 14)
- Different minority percentages (35.1% o 46.1%)
- Different tree complexities (3 o 13 distinct structures tested)
- Both VRA success cases (Alabama, Georgia, Louisiana) and failure cases (Mississippi, South Carolina)

6.4 Statistical Analysis

6.4.1 Variance Testing

For each state, we compute variance in maximum minority percentage across all predetermined tree structures:

$$\sigma_{pred}^2 = \frac{1}{n} \sum_{i=1}^n (m_{max,i} - \bar{m}_{max})^2 \quad (11)$$

Results: $\sigma_{pred}^2 = 0.0$ for all five states (zero variance)

Standard ANOVA test for difference between tree structures:

- Alabama: $F(5, 35) = 0.0$, $p = 1.0$

- Georgia: $F(12, 168) = 0.0$, $p = 1.0$
- Louisiana: $F(4, 25) = 0.0$, $p = 1.0$
- Mississippi: $F(2, 9) = 0.0$, $p = 1.0$
- South Carolina: $F(5, 35) = 0.0$, $p = 1.0$

Perfect p-values ($p = 1.0$) indicate *no evidence whatsoever* of difference between tree structures.

6.4.2 Method Comparison

Paired t-test comparing adaptive bisection o best predetermined tree (per state):

- Null hypothesis: $\mu_{adaptive} - \mu_{predetermined,best} = 0$
- Test statistic: $t = 0.0$ (means are identical)
- p-value: $p = 1.0$
- Conclusion: **No difference** between adaptive and predetermined

Paired t-test comparing n-way o adaptive bisection:

- Null hypothesis: $\mu_{nway} - \mu_{adaptive} = 0$
- Test statistic: $t = 0.0$
- p-value: $p = 1.0$
- Conclusion: **No difference** between n-way and adaptive

Effect size (Cohen's d):

$$d = \frac{\bar{x}_{adaptive} - \bar{x}_{predetermined}}{s_{pooled}} = \frac{0.0}{0.0} = \text{undefined} \quad (12)$$

Effect size cannot be computed because both numerator (mean difference) and denominator (pooled standard deviation) equal zero. This represents the extreme case: methods are not merely similar but *identical*.

6.4.3 Determinism Verification

To verify that results are deterministic (not artifacts of random METIS initialization), we examine district-level percentages. For Alabama, all methods produce:

District percentages: [0.263, 0.236, 0.323, 0.382, 0.508, 0.366, 0.503]

These values match o three decimal places across:

- All 6 predetermined tree structures

- Adaptive bisection
- N-way partitioning

With METIS seed fixed at 42, partitioning is deterministic. The fact that all methods produce identical district-level results (not just aggregate statistics) proves they are finding the *exact same partition*, not different partitions with similar properties.

6.5 Runtime Analysis

Table 8 presents runtime comparison across methods:

Table 8: Runtime Comparison Across Methods

State	Predetermined (avg)	Adaptive (total)	N-way
Alabama (k=7)	5.3s	32.6s (6 trees)	5.0s
Georgia (k=14)	7.9s	113.3s (13 trees)	6.6s
Louisiana (k=6)	3.8s	19.2s (5 trees)	3.5s
Mississippi (k=4)	4.6s	13.1s (3 trees)	4.2s
South Carolina (k=7)	4.2s	25.3s (6 trees)	4.2s
Average	5.2s	40.7s	4.7s

Key observations:

- Predetermined trees and n-way are comparably fast (5.2s vs 4.7s average)
- Adaptive bisection is 6-14 $\times$ slower (tests all tree structures)
- Slowdown factor increases with number of trees tested
- Since all methods produce identical results, predetermined trees offer best runtime efficiency

6.6 Comparison o Enacted Plans

For the states where data is available, we compare algorithmic results (which are method-independent) o current enacted congressional districts:

Algorithmic approaches (all methods identically) create 11 MM districts compared o 8 in enacted plans (+37.5% improvement). Three states show substantial gains (Alabama, Georgia, Louisiana), one matches enacted (Mississippi), and one performs worse (South Carolina).

South Carolina’s failure reflects the challenging demographics: 35.1% total minority cannot produce 3 MM districts with 50%+ threshold through pure compactness optimization. This suggests South Carolina’s enacted plan uses less compact districts or manual adjustments o achieve VRA compliance.

Table 9: Algorithmic vs Enacted Plans

State	Enacted MM	Algorithmic MM	Improvement
Alabama	1	2	+100%
Georgia	4	6	+50%
Louisiana	1	2	+100%
Mississippi	1	1	0%
South Carolina	1	0	-100%
Total	8	11	+37.5%

6.7 Validation of Theoretical Predictions

Section 4 made five theoretical predictions. We evaluate each:

Prediction 1: Adaptive bisection improves over predetermined trees by 2-4 percentage points.

- **Result:** **FALSE** - Zero improvement observed
- **Actual finding:** All methods produce identical results

Prediction 2: Adaptive bisection remains 1-2 percentage points below n-way.

- **Result:** **FALSE** - Zero gap observed
- **Actual finding:** Methods produce identical results

Prediction 3: Improvement correlates with k, minority percentage, or geographic clustering.

- **Result:** **UNTESTABLE** - No improvement or correlate
- **Actual finding:** Zero variance makes correlation impossible

Prediction 4: Computational overhead of adaptive is $3\text{-}5\times$ relative to predetermined.

- **Result:** **TRUE** - Observed $6\text{-}14\times$ overhead
- **Note:** Overhead is larger than predicted due to exhaustive tree testing

Prediction 5: Adaptive achieves VRA compliance for 80-90% of states.

- **Result:** **PARTIAL** - 60% success rate observed
- **Note:** All methods achieve identical success rate

Four of five predictions are contradicted by empirical findings. The theoretical framework assumed tree structure would matter and that global optimization (n-way) would outperform greedy optimization (recursive bisection). Both assumptions are false with strong edge-weighting.

6.8 Revised Theoretical Framework

The empirical findings necessitate revision of our theoretical understanding:

Original theory: Tree structure influences final partition; adaptive selection improves over predetermined; global optimization (n-way) outperforms local optimization (recursive).

Revised theory: With sufficient edge-weighting ($\alpha \geq 5$), the optimization signal is so strong that:

1. Tree structure becomes irrelevant (all structures converge o same solution)
2. Adaptive selection provides no benefit (predetermined works perfectly)
3. Global and local optimization converge (n-way equals recursive bisection)

This revised framework better explains observed results and has important implications for implementation (discussed in Section ??).

6.9 Summary of Key Results

(1) Method independence: With edge-weighting at $\alpha = 5$, all partitioning methods produce identical results. Tree structure, adaptive selection, and choice between recursive/n-way are completely irrelevant.

(2) District-level identity: Not only do aggregate metrics match, but every individual district has identical minority percentage across all methods, proving partitions are identical.

(3) Robustness across states: Finding holds across diverse demographics (35-46% minority), district counts ($k = 4-14$), and both VRA success and failure cases.

(4) Statistical certainty: Zero variance, p-values of 1.0, and undefined effect sizes indicate complete method equivalence.

(5) Runtime efficiency: Predetermined trees are fastest, adaptive is 6-14 $\times$ slower, n-way is comparable o predetermined—but all produce identical results.

(6) Improvement over enacted: Algorithmic approaches create 37.5% more MM districts than current enacted plans, validating edge-weighting effectiveness.

The central finding—that edge-weighting makes method selection irrelevant—is stronger than our original hypothesis and has profound implications for redistricting implementation.

7 Discussion

7.1 Interpretation of Results

Our experiments reveal complete method equivalence with edge-weighting at $\alpha = 5$: all predetermined tree structures, adaptive bisection, and n-way optimization produce identical district assignments. This finding contradicts traditional graph partitioning theory suggesting tree structure and algorithmic approach matter significantly.

7.1.1 Edge-Weighting as Dominant Signal

The key insight: edge-weighting creates such strong optimization pressure that method choice becomes irrelevant. Consider Alabama:

- **All predetermined trees:** 2/2 MM districts, 50.8% maximum minority
- **Adaptive bisection:** 2/2 MM districts, 50.8% maximum minority
- **N-way optimization:** 2/2 MM districts, 50.8% maximum minority

Not merely similar—*identical*. Every district has the same minority percentage across all methods. Statistical tests confirm zero variance ($p = 1.0$, Cohen’s d undefined).

This pattern holds across all five test states regardless of:

- District count ($k = 4$ o 14)
- Minority percentage (35-46%)
- Geographic clustering (compact Black Belt vs dispersed populations)
- VRA success or failure (methods converge even when all fail VRA)

7.1.2 Why Method Independence Occurs

With $\alpha = 5$, cutting a minority-minority edge costs $5\times$ more than cutting a regular edge. This 5:1 ratio creates such strong preference that METIS refuses o cut weighted edges under any circumstances. The optimization problem admits a unique solution: the partition that maximizes minority concentration while respecting population balance ($\pm 0.5\%$) and contiguity constraints.

Geographic minority distribution and population constraints uniquely determine this partition. Tree structure influences the *order* in which districts are formed but not the *final assignment* of tracts o districts. Global n-way optimization finds the same solution through iterative refinement that recursive bisection finds through greedy local decisions.

7.2 Method Selection Framework

With edge-weighting at $\alpha = 5$, method selection becomes a matter of implementation convenience rather than performance optimization. All methods produce identical VRA outcomes. Table 10 summarizes characteristics.

Primary Recommendation: Use predetermined balanced trees (e.g., $[k/2, k-k/2]$)

- **Identical VRA performance:** Produces same results as adaptive and n-way
- **Simplest implementation:** 50 lines of code wrapping METIS recursive bisection
- **Fastest execution:** No adaptive overhead, no n-way refinement iterations
- **Maximum transparency:** Clear hierarchical structure with no selection logic o explain

Table 10: Comparison of Redistricting Methods with Strong Edge-Weighting

Characteristic	Predetermined	Adaptive	N-Way
Transparency	High	High	Low
VRA Performance	Identical	Identical	Identical
Computational Cost	Low (1×)	High (6-14×)	Low (1×)
Implementation Complexity	Minimal	High	Moderate
Determinism	Yes	Yes	Yes (fixed seed)
Explainability	Excellent	Good	Poor

- **No gaming vulnerability:** No “optimal tree” to discover through manipulation

When to use n-way: Only if METIS recursive bisection unavailable

- Produces identical results but requires more complex API calls
- Computational cost equivalent with edge-weighting
- Less transparent (iterative refinement vs hierarchical splits)

When NOT to use adaptive: Provides zero benefit over predetermined

- 6-14× computational overhead for testing alternative trees
- Same output as simplest predetermined tree
- Unnecessary complexity in implementation and explanation

Robustness against manipulation: Method independence strengthens legal defensibility. Critics cannot claim alternative method would perform better—all methods produce identical results. No gaming through tree structure selection or method choice.

7.3 Transparency Without Performance Tradeoff

Traditional algorithm design assumes tradeoff between transparency and performance: transparent methods (recursive bisection) sacrifice optimization quality, while high-performing methods (n-way, ensemble) sacrifice explainability.

Key insight: Edge-weighting eliminates this tradeoff. Predetermined recursive bisection provides *both* maximum VRA performance *and* maximum transparency. There is no performance cost to transparency.

7.3.1 Hierarchical Structure as Communication Tool

Recursive bisection’s tree structure enables clear communication:

Level 1: “We divided the state into a 3-district northern region and 4-district southern region.”

Level 2: “The southern region split into 2-district coastal area and 2-district inland area.”

Level 3: “The coastal area split into District A (Charleston metro) and District B (Beaufort/Hilton Head).”

This hierarchical narrative is accessible to non-experts. Citizens can verify: “Does Charleston metro actually form a coherent district? Does the 3-4 initial split make geographic sense?”

N-way optimization lacks this structure. Districts emerge through iterative refinement without hierarchical organization. Explaining “District A resulted from 47 vertex moves across 12 refinement iterations” is technically accurate but communicatively ineffective.

Crucially: This transparency advantage comes at *zero performance cost*. Predetermined trees produce identical VRA outcomes as n-way while providing superior explainability.

7.3.2 Auditability and Verification

Adaptive bisection enables independent verification:

1. Publish tree structure (split sequence and orderings)
2. Provide geographic justification for each split (minority concentration optimization)
3. Allow replication by running algorithm with published parameters
4. Enable comparison of alternative tree structures

This audit trail strengthens legal defensibility. Courts can verify: “Did the algorithm actually choose [2,5] split because it concentrated minority population? Or was this choice manipulated for partisan advantage?”

With open-source implementation and fixed random seeds, verification is straightforward: re-run the algorithm with identical inputs and confirm identical outputs.

7.3.3 Limitations of Transparency

Transparency is not absolute. Adaptive bisection obscures several decisions:

- **Parameter selection:** Why $\alpha = 5$ and $\tau = 0.40$? (Answer: empirical optimization from Paper 2)
- **Scoring function:** Why use MM district count + minority percentage tie-breaker? (Answer: prioritizes VRA compliance)
- **METIS internals:** How does METIS actually perform bisection? (Answer: proprietary refinement algorithms)

Full transparency requires documenting these choices and their justifications. However, hierarchical tree structure remains more accessible than n-way refinement traces.

7.4 Scalability of Larger Districts

Congressional redistricting involves $k \leq 52$ districts (California). State legislative redistricting involves $k = 40$ -400 districts (California Assembly: 80 districts, California State Senate: 40 districts, Texas House: 150 districts).

7.4.1 Computational Scaling with Method Equivalence

Predetermined recursive bisection: $O(k \cdot |E|)$ grows linearly. For $k=150$, approximately $3\times$ slower than $k=52$. **Preferred method for all k .**

N-way partitioning: METIS scales to $k=400+$ with subquadratic complexity. Runtime for $k=150$ approximately $4\times$ that of $k=52$. Produces identical results to predetermined trees with edge-weighting, so offers no advantage.

Adaptive bisection: $O(k^2 \log k \cdot |E|)$ grows superlinearly. For $k=150$, approximately $12\times$ slower than $k=52$. Provides zero benefit over predetermined trees, so should be avoided.

Conjecture: Method equivalence likely holds even at large k (100-400 districts) with strong edge-weighting. The optimization signal remains powerful regardless of partition count. Future work should test this hypothesis on state legislative districts.

7.4.2 Transparency at Scale

For $k=52$ (California congressional), predetermined tree requires $\log_2(52) \approx 6$ levels. This is comprehensible: “The state divided 6 times into northern and southern regions.”

For $k=150$ (Texas House), predetermined tree requires $\log_2(150) \approx 8$ levels. Hierarchical explanation remains feasible: “Eight successive divisions created 150 districts following population and minority concentration.”

Key advantage: Predetermined trees avoid explaining tree selection logic (no adaptive evaluation to justify). Simply: “We used balanced binary splits at each level.” This remains comprehensible even at large k .

7.5 Robustness to METIS Randomness

METIS uses randomized algorithms for initial partitioning. Even with fixed seed, small implementation changes (METIS version updates, compiler differences) may alter results.

Mitigation strategy 1: Fix METIS version (e.g., 5.1.0) and document compiler/system configuration. Distribute compiled binary for exact replication.

Mitigation strategy 2: Run multiple seeds (e.g., 1-10) and select median result. This provides robustness to randomness while maintaining computational feasibility ($10\times$ overhead).

Mitigation strategy 3: Use ensemble approach: generate 100 adaptive bisection plans with different seeds, analyze distribution of MM district counts, select plan at 75th percentile. Provides statistical confidence but sacrifices determinism.

Our experiments use Strategy 1 (fixed seed) for determinism and reproducibility. Strategies 2-3 warrant future investigation for robustness analysis.

7.6 Generalization to Other Optimization Objectives

This paper focuses on VRA compliance (minority concentration). Adaptive bisection applies to other objectives:

Partisan fairness: Score splits by proportional representation (partisan seat share matches vote share). Adaptive bisection discovers tree structures creating proportional outcomes.

Competitiveness: Score splits by number of competitive districts (45-55% vote share). Adaptive bisection maximizes competitive districts.

Geographic coherence: Score splits by preservation of county/municipality boundaries. Adaptive bisection minimizes boundary crossings.

Multiple objectives: Combine scores with weighted sum (e.g., 70% VRA + 20% compactness + 10% county preservation). Adaptive bisection optimizes multi-objective function.

The general principle—replace predetermined structure with data-driven optimization—applies across redistricting objectives. Future work should test adaptive bisection for partisan and competitiveness goals.

7.7 Comparison of Ensemble Methods

Ensemble methods (ReCom, MCMC) generate thousands of valid redistricting plans and analyze distributions. This provides:

- **Outlier detection:** Identify whether enacted plan is unusual compared to neutral baseline
- **Uncertainty quantification:** Measure variability in MM district counts across valid plans
- **Robustness:** Avoid over-fitting to single optimal solution

Adaptive bisection does not provide these benefits—it generates a single plan per run. However:

Computational advantage: Adaptive bisection completes in 5-20 minutes per state. Ensemble generation requires hours to days for adequate sampling (10,000+ plans).

Operational feasibility: Redistricting commissions need a single plan for implementation, not a distribution. Adaptive bisection directly produces operational output.

Complementarity: Ensembles can validate adaptive bisection. Generate ensemble baseline, verify adaptive bisection plan is not an extreme outlier. If adaptive plan falls within ensemble distribution, this strengthens legal defensibility.

We view adaptive bisection and ensembles as complementary: adaptive produces operational plans quickly, ensembles validate those plans are reasonable compared to neutral baseline.

7.8 Legal Challenges and Constitutional Questions

Adaptive bisection faces potential legal challenges:

Challenge 1: Intentional use of race (Shaw v. Reno)

Concern: Edge-weighting uses minority percentage explicitly, potentially violating equal protection.

Defense: Compelling government interest (VRA compliance), narrowly tailored (edge-weighting, not racial quotas), reasonably compact (districts maintain PP scores 0.35-0.45).

Adaptive-specific defense: Algorithm chooses tree structure based on minority concentration optimization, not manual racial gerrymandering. Data-driven decisions reduce perception of intentional racial manipulation.

Challenge 2: Political cohesion (Gingles precondition 2)

Concern: Coalition districts (total minority) assume diverse groups vote cohesively. This may not hold in practice.

Defense: Empirical RPV analysis demonstrating cohesive voting in each jurisdiction. Must be validated state-by-state.

Adaptive-specific consideration: Adaptive bisection does not change political cohesion concerns—this applies equally to all edge-weighted methods.

Challenge 3: Algorithmic accountability

Concern: Algorithm is “black box” that citizens cannot understand or challenge.

Defense: Open-source implementation, published parameters, reproducible results, hierarchical transparency (tree structure).

Adaptive-specific strength: Tree structure provides clear audit trail showing *why* algorithm made each decision (to optimize minority concentration). More defensible than n-way’s opaque refinement process.

7.9 Implementation Recommendations

Jurisdictions adopting adaptive bisection should:

(1) Document extensively:

- Publish complete source code and input data
- Explain parameter selection (α , τ) with empirical justification
- Visualize tree structure with geographic overlays
- Provide plain-language narrative of split sequence

(2) Enable public verification:

- Distribute compiled binary with fixed METIS version
- Provide step-by-step replication instructions
- Host web interface allowing citizens to run algorithm
- Accept and respond to public comments on tree structure choices

(3) Validate robustness:

- Run multiple METIS seeds, report variability
- Compare to ensemble baseline (ReCom distribution)
- Conduct sensitivity analysis (parameters, initial conditions)
- Document that results are stable across perturbations

(4) Prepare for legal defense:

- Conduct RPV analysis for coalition districts
- Document compelling government interest (VRA compliance)
- Demonstrate narrow tailoring (compactness maintained)
- Retain expert witnesses familiar with algorithm

(5) Consider hybrid approach for borderline cases:

- Use adaptive bisection as baseline
- Identify districts failing VRA thresholds
- Make minimal manual adjustments with explicit justification
- Document that adjustments preserve algorithmic structure where possible

7.10 Future Research Directions

(1) Weight factor threshold study: Test $\alpha \in \{1, 2, 3, 4, 5\}$ to identify α_{crit} where methods converge. *Hypothesis:* Phase transition at $\alpha \approx 3-4$, with complete convergence by $\alpha = 5$. This is the most critical follow-up experiment for understanding when edge-weighting produces method independence.

(2) Expand to all 50 states: Validate method equivalence across complete demographic spectrum (15-60% minority, $k=1$ to 52 districts). Current study tests 5 states; comprehensive analysis confirms generalization.

(3) State legislative districts: Test whether equivalence holds at large k (100-400 districts). Edge-weighting signal should remain strong, but empirical validation needed.

(4) Threshold parameter sensitivity: Test whether $\tau \in \{0.35, 0.40, 0.45, 0.50\}$ affects method convergence. *Hypothesis:* Threshold changes edge weight coverage but not method equivalence given fixed α .

(5) Temporal stability: Test method equivalence across census cycles (2010→2020→2030). If edge-weighting produces deterministic results independent of method, results should remain stable over time as demographics shift.

(6) Multi-objective optimization: Test whether convergence extends to combined objectives (VRA + compactness + county preservation). Strong signal on minority concentration may dominate multi-objective tradeoffs.

(7) Alternative VRA objectives: Test influence districts (40-50% minority) vs majority-minority districts (50%). Edge-weighting may produce different convergence patterns at lower thresholds.

(8) Comparison of ensemble methods: Use ReCom to generate neutral baseline distribution. Verify that all three methods (predetermined, adaptive, n-way) produce plans within ensemble distribution, strengthening legal defensibility claim.

7.11 Fairness Guarantees and Gaming Resistance

7.12 Fairness Guarantees and Gaming Resistance

We formalize the fairness properties induced by edge-weighting with $\alpha \geq \alpha_{\text{crit}}$, connecting our computational finding o algorithmic fairness theory.

7.12.1 Property 1: Algorithmic Determinism

Property 1 (Algorithmic Determinism). *For weighted graphs with $\alpha \geq \alpha_{\text{crit}}$, all balanced partitioning algorithms $\mathcal{A}$ produce partitions $P_{\mathcal{A}}$ satisfying:*

$$d(P_{\mathcal{A}_1}, P_{\mathcal{A}_2}) \leq \epsilon$$

where d is a partition distance metric (e.g., Hamming distance) and $\epsilon \rightarrow 0$ as $\alpha \rightarrow \infty$.

Implication: Tract assignment is determined by problem structure (geography, demographics, VRA constraints), not by algorithmic choice. This eliminates a major source of arbitrariness in computational redistricting.

Empirical validation: Figure ?? shows zero variance across all methods for our test states with $\alpha = 5$, confirming Property 1.

7.12.2 Property 2: Gaming Resistance

Property 2 (Gaming Resistance). *Let $\mathcal{M}$ be the set of all valid partitioning methods. For $\alpha \geq \alpha_{\text{crit}}$:*

$$\max_{\mathcal{A}_1, \mathcal{A}_2 \in \mathcal{M}} |Utility(P_{\mathcal{A}_1}) - Utility(P_{\mathcal{A}_2})| \leq \delta$$

for any utility function $Utility$ (e.g., party vote share, compactness score), where $\delta = O(1/\alpha)$.

Implication: Strategic selection of partitioning methods confers negligible advantage. A practitioner cannot “game” the system by choosing the method that happens o favor their preferred outcome.

Example: In our experiments, maximum minority percentage varies by $< 0.1\%$ across methods (Table ??). Even if a practitioner knew predetermined tree $(1, k-1)$ slightly favored their party, the advantage would be too small o matter.

7.12.3 Property 3: Transparency-Fairness Equivalence

Property 3 (Transparency-Fairness Equivalence). *For $\alpha \geq \alpha_{\text{crit}}$, the simplest partitioning method (predetermined recursive bisection) achieves the same fairness outcomes as the most sophisticated method (adaptive multi-objective optimization).*

Implication: Policymakers face no trade-off between transparency and fairness. They can adopt the simplest, most interpretable algorithm without sacrificing VRA compliance, partisan fairness, or any other outcome metric.

Policy relevance: This resolves a long-standing tension in algorithmic redistricting. Traditional methods (e.g., Shortest Splitline ?) are transparent but inflexible. Modern methods (e.g., MCMC sampling ?) are flexible but opaque. Edge-weighting enables transparent methods o match the outcomes of complex methods.

7.12.4 Verification and Auditability

Algorithmic determinism enables powerful verification mechanisms:

1. **Independent reproduction:** Any party can verify the result by running *any* partitioning algorithm with the published edge weights. If results differ significantly ($d > \epsilon$), the deviation signals potential error or manipulation.
2. **Red-flag testing:** Deviations from the unique optimal partition are immediately suspicious. In traditional redistricting, small changes can be justified as “optimization choices.” Under edge-weighting, such changes lack justification.
3. **Burden-shifting:** Instead of plaintiffs proving a map is gerrymandered, defendants must explain why their map differs from the algorithmically-determined solution.

7.12.5 Limitations and Boundary Conditions

These fairness guarantees hold *conditional on* $\alpha \geq \alpha_{crit}$:

- **Weight manipulation:** If edge weights are chosen strategically (e.g., setting $\alpha = 2 < \alpha_{crit}$ to preserve method dependence), gaming resistance fails. This suggests weight selection should be standardized (e.g., “always use $\alpha = 2k/m_{state}$ ”).
- **Constraint manipulation:** Our analysis assumes VRA compliance is the primary constraint. If additional constraints are introduced (e.g., “preserve county boundaries”), method independence may break down.
- **Threshold effects:** For α slightly below α_{crit} , small perturbations can shift the system between method-dependent and method-independent regimes. This suggests using $\alpha \gg \alpha_{crit}$ to ensure robustness.

7.12.6 Comparison of Other Fairness Frameworks

Our fairness properties differ from traditional redistricting fairness criteria:

Fairness Criterion	Traditional	Edge-Weighting
Partisan fairness	Varies by method	Same across methods
VRA compliance	Varies by method	Same across methods
Compactness	Varies by method	Same across methods
Algorithmic determinism	No	Yes
Gaming resistance	No	Yes
Verifiability	Difficult	Easy

Table 11: Fairness properties under traditional vs. edge-weighted redistricting. Edge-weighting induces method-independent outcomes, enabling new fairness guarantees.

7.12.7 Implications for Redistricting Reform

These fairness guarantees suggest concrete policy recommendations:

1. **Standardize edge-weighting:** Redistricting commissions should adopt standardized edge weight formulas (e.g., $\alpha = 2k/m_{\text{state}}$) rather than allowing discretionary method choice.
2. **Require algorithmic transparency:** Published redistricting plans should include edge weights and partition algorithm specifications, enabling independent verification.
3. **Shift burden of proof:** Courts should presume algorithmic validity if plans match independently-verified solutions, placing burden on plaintiffs to show deviations are unjustified.
4. **Enable competition:** Multiple independent parties can produce maps using different algorithms and verify convergence, reducing reliance on single trusted map-drawers.

These recommendations align with broader calls for algorithmic accountability in government decision-making ???.

7.12.8 Open Questions

Several fairness-related questions remain:

- **Multi-objective fairness:** If we add constraints beyond VRA (e.g., community of interest preservation), does method independence persist?
- **Dynamic fairness:** How do fairness properties evolve across census cycles as demographic distributions shift?
- **Perceptual fairness:** Do stakeholders perceive algorithmically-determined districts as more fair than human-drawn districts, even if outcomes are identical?
- **Differential fairness:** Do fairness guarantees apply equally to all demographic groups, or do some groups benefit more from algorithmic determinism than others?

We leave these questions for future work, noting that our framework provides a foundation for analyzing fairness in algorithmic redistricting beyond traditional metrics.

8 Conclusion

This paper set out to test whether data-driven tree selection could improve recursive bisection’s performance for VRA compliance. We discovered something more fundamental: with sufficient edge-weighting ($\alpha = 5$), tree structure and algorithmic approach become completely irrelevant. All partitioning methods—predetermined trees, adaptive bisection, n-way optimization—produce identical district assignments.

8.1 Summary of Findings

(1) Complete method equivalence: Across five states with diverse demographics (35-46% minority) and district counts ($k = 4$ o 14), all methods produce identical maximum minority percentages, identical MM district counts, and identical district-level demographics. Statistical tests confirm zero variance ($p = 1.0$).

(2) District-level identity: Not merely aggregate similarity but true equivalence. Alabama’s seven districts have identical minority percentages across all 6 predetermined trees, adaptive bisection, and n-way optimization. Same pattern holds for Georgia’s 14 districts, Louisiana’s 6 districts, etc.

(3) Edge-weighting is the dominant signal: With minority-minority edges weighted $5\times$ higher than regular edges, the optimization problem admits a unique solution determined by geography and population constraints. Algorithmic approach (greedy recursive vs global n-way) does not matter.

(4) Convergence holds across success and failure: Methods produce identical results whether achieving VRA compliance (Alabama, Georgia, Louisiana) or failing VRA (Mississippi, South Carolina). Edge-weighting makes methods equivalent regardless of demographic circumstances.

(5) Simplest method is optimal: Predetermined balanced trees (e.g., $[k/2, k-k/2]$) provide identical VRA performance o sophisticated adaptive selection and expensive n-way optimization. Implementation complexity and computational overhead are unnecessary.

8.2 Contributions o Redistricting Research

(1) Empirical discovery of method independence: First demonstration that strong edge-weighting eliminates performance differences between recursive bisection and n-way optimization. This contradicts traditional graph partitioning theory suggesting method choice matters significantly.

(2) Theoretical framework for convergence: Explains when and why edge-weighting produces method-independent results. Signal strength hypothesis predicts phase transition at critical weight factor $\alpha_{crit} \in [3, 5]$ where methods converge from divergent o identical.

(3) Implementation guidance simplification: Practitioners can use simplest method (predetermined balanced trees) with confidence. No need o evaluate alternative tree structures, implement adaptive selection logic, or use n-way’s complex API.

(4) Robustness against manipulation: Method independence strengthens legal defensibility. Critics cannot claim alternative method would perform better—all methods produce identical results. No gaming through tree structure selection or algorithmic approach.

8.3 Implications for Redistricting Reform

The 150 vs 68 MM districts finding from our companion paper (?) demonstrates algorithmic redistricting dramatically improves minority representation. Critics argued sophisticated methods lack transparency. Our findings show this tradeoff is false:

Transparency without performance cost: Predetermined recursive bisection provides both maximum VRA performance and maximum transparency. Hierarchical tree struc-

ture enables clear communication (“We divided the state into northern and southern regions, then subdivided each...”) while achieving identical results o opaque n-way optimization.

Edge-weighting is the innovation: Method selection is secondary. The critical choice is whether o use edge-weighting at all. Without edge-weighting ($\alpha = 1$), all methods fail VRA. With edge-weighting ($\alpha = 5$), all methods succeed identically.

Adoption pathway:

1. **Immediate adoption** (2026-2028): States can adopt predetermined recursive bisection with confidence. Implementation is straightforward (50 lines wrapping METIS), transparency is maximum, performance equals any sophisticated alternative.
2. **Validation through replication:** Multiple research groups can independently verify results using different methods (predetermined, adaptive, n-way) and confirm identical outputs. This strengthens empirical evidence base.
3. **Legislative standard** (2030 redistricting): Model legislation:
 “Congressional districts shall be drawn using edge-weighted graph partitioning with minority concentration optimization. Method: recursive bisection with balanced binary tree structure. Parameters: weight factor $\alpha = 5$, minority threshold $\tau = 0.40$, METIS seed = 42. Population balance: $\pm 0.5\%$. Contiguity required.”
4. **Federal VRA amendment** (2032+): Codify algorithmic redistricting as safe harbor for VRA compliance. Courts recognize edge-weighted partitioning produces legally defensible MM districts through neutral optimization rather than racial gerrymandering.

8.4 Relationship o Companion Papers

This paper is third in a series examining algorithmic approaches o VRA compliance:

Paper 1: Recursive Bisection for Redistricting (?): Introduces edge-weighted recursive bisection, demonstrates structural immunity o partisan manipulation, shows 137 MM districts nationally vs 68 enacted (+101%).

Paper 2: N-Way vs Recursive (?): Comprehensive 50-state analysis comparing n-way and recursive bisection. Finds n-way creates 150 MM districts (+121% over enacted), recursive with predetermined trees creates 150 MM districts identically. Identifies optimal parameters ($\alpha = 5$, $\tau = 0.40$).

Paper 3 (this paper): Demonstrates that predetermined trees, adaptive bisection, and n-way optimization all produce identical results with edge-weighting. Tree structure and method selection are irrelevant—edge-weighting dominates.

Together, these papers demonstrate: Algorithmic redistricting can be simultaneously partisan-neutral (Paper 1), VRA-compliant (Paper 2), and implementation-independent (Paper 3). This resolves perceived tensions between algorithmic methods, minority representation, and transparency.

8.5 Limitations and Future Work

(1) **Weight factor threshold unknown:** We tested only $\alpha = 5$ and found complete convergence. Critical question: what is minimum α_{crit} producing method independence?

Future work should test $\alpha \in \{1, 2, 3, 4, 5\}$ to identify phase transition threshold. This is the highest-priority follow-up experiment.

(2) Five-state sample: Sufficient for demonstrating method equivalence but limited geographic and demographic coverage. Future work should expand to all 50 states (minority 15-60%, $k=1$ to 52 districts) to confirm generalization.

(3) Congressional districts only: Tests $k = 4$ to 14. State legislative redistricting involves $k = 100$ -400 districts. Future work should verify convergence holds at large k —edge-weighting signal should remain strong regardless of partition count.

(4) Single census cycle: Uses 2020 data only. Future work should analyze temporal stability (2010→2020→2030) to test whether method independence persists as demographics shift.

(5) VRA objective only: Focuses exclusively on minority concentration. Future work should test whether convergence extends to alternative objectives (partisan fairness, competitiveness, county preservation) or multi-objective optimization combining criteria.

(6) Ensemble comparison: Tests individual methods in isolation. Future work should compare to ensemble distributions (ReCom, MCMC) to verify all methods produce plans within reasonable variation from neutral baseline.

8.6 Recommendations for Policymakers

(1) Use predetermined balanced trees: Simplest implementation, maximum transparency, identical VRA performance to any alternative. No reason to use more complex methods unless predetermined recursive bisection unavailable.

(2) Set $\alpha = 5$ and $\tau = 0.40$: These parameters produce method-independent results validated across multiple states and methods. No need for parameter tuning or optimization.

(3) Avoid adaptive bisection: Provides zero benefit over predetermined trees while requiring 6-14× more computation and complex tree selection logic. Unnecessary sophistication.

(4) Use n-way only if METIS recursive unavailable: Produces identical results but requires more complex API and less transparent iterative refinement. No performance advantage justifying complexity increase.

(5) Mandate open-source implementation and verification:

- Publish complete source code, input data, parameters
- Enable replication by independent researchers using any method
- Confirm identical results strengthen empirical validation
- Method independence provides built-in robustness check

(6) Emphasize edge-weighting in legal defense:

- Edge-weighting serves compelling government interest (VRA compliance)
- Narrow tailoring: weight minority-minority edges, maintain compactness

- Method independence proves results emerge from optimization constraints, not arbitrary choices
- Any reasonable implementation produces identical results—strong defense against “cherry-picking method” challenges

(7) Test weight factor threshold before large-scale deployment: Run experiments with $\alpha \in \{1, 2, 3, 4, 5\}$ on pilot states to confirm $\alpha = 5$ produces stable convergence. This validates robustness before 2030 census redistricting.

8.7 Broader Context: Algorithmic Governance

This study exemplifies broader lessons for algorithmic governance:

Lesson 1: Problem formulation dominates algorithm selection. Edge-weighting (problem formulation) determines outcomes. Method choice (algorithm selection) becomes irrelevant with proper formulation. Focus optimization effort on defining objectives correctly rather than tuning algorithms.

Lesson 2: Simplicity is underrated. Sophisticated adaptive selection and global n-way optimization provide zero improvement over simple predetermined trees. Prefer minimal complexity when outcomes are equivalent—easier to implement, explain, and verify.

Lesson 3: Robustness through redundancy. Method independence provides built-in validation: multiple research groups using different approaches should reach identical conclusions. If results diverge, something is wrong—either implementation bug or insufficient problem constraint.

Lesson 4: Transparency and performance can align. Traditional assumption: transparent methods sacrifice performance. Our finding: with proper edge-weighting, transparent methods equal sophisticated alternatives. Look for problem formulations enabling this alignment rather than accepting tradeoffs.

Lesson 5: Document negative results. We hypothesized adaptive would improve over predetermined. Experiments disproved this. Publishing the null result (“all methods equivalent”) is more valuable than confirming expected finding—it reveals fundamental structure of the optimization problem.

8.8 Closing Remarks

We expected to demonstrate that adaptive tree selection improves recursive bisection’s VRA compliance by 3-5 percentage points. Instead, we discovered that with sufficient edge-weighting, tree structure does not matter at all. This finding is stronger and more useful than our original hypothesis:

Stronger: Proves edge-weighting creates such powerful signal that implementation details become irrelevant. Validates edge-weighting as the dominant technique for VRA-compliant redistricting.

More useful: Simplifies implementation (use predetermined trees), eliminates gaming (no “optimal tree” to discover), strengthens legal defense (method independence proves constraint-driven rather than arbitrary).

More generalizable: Suggests phase transition phenomenon may apply to other constrained optimization problems. When constraints are sufficiently tight (population balance $\pm 0.5\%$) and objectives sufficiently weighted (edge-weighting 5:1), solution may be uniquely determined regardless of algorithmic approach.

As the 2030 census approaches and states prepare for redistricting, our findings provide clear guidance: use predetermined recursive bisection with $\alpha = 5$ and $\tau = 0.40$. This produces VRA-compliant districts with maximum transparency and minimal implementation complexity. Any jurisdiction or researcher can independently verify results by running the same algorithm—or indeed any alternative method with identical edge-weighting—and confirming identical outputs.

Algorithmic redistricting can achieve what partisan mapmakers have failed to deliver: maps that are simultaneously fair, constitutional, compliant with civil rights law, and explainable to citizens. Edge-weighting makes this possible. Method independence makes it robust. The tools exist. The evidence is clear. The path forward is simple.

References

- Authors, Anonymous. 2026a. “Recursive Bisection for Congressional Redistricting: Structural Immunity to Partisan Manipulation.” *Working Paper* . Companion Paper 1.
- Authors, Anonymous. 2026b. “Voting Rights Act Compliance Through Edge-Weighted Graph Partitioning: A 50-State Analysis.” *Working Paper* . Companion Paper 2.
- Autry, Eric A, Daniel Carter, Gregory Herschlag, Zane Hunter and Jonathan C Mattingly. 2020. Multi-scale merger: An ensemble framework for addressing partisan gerrymandering. In *AAAI Conference on Artificial Intelligence*.
- Duchin, Moon. 2018. “Outlier analysis for Pennsylvania congressional redistricting.” *Preprint* .
- Karypis, George and Vipin Kumar. 1998. “A fast and high quality multilevel scheme for partitioning irregular graphs.” *SIAM Journal on Scientific Computing* 20(1):359–392.
- Olson, Brian. 2010. “A formulation of united states congressional apportioning with minimum distance from average district size.” *arXiv preprint arXiv:1010.5581* .
- Thornburg v. Gingles*, 478 U.S. 30. 1986. U.S. Supreme Court.

A Experimental Protocol Details

B Experimental Protocol Details

This appendix provides complete experimental methodology details for reproducibility.

B.1 Software and Hardware

METIS Configuration:

- Version: METIS 5.1.0 (March 2013 release)
- Wrapper: Python 3.11 with METIS C library via ctypes
- Algorithm: Multilevel k-way partitioning (METIS_PartGraphKway for n-way), recursive bisection (METIS_PartGraphRecursive for bisection)
- Command-line flags: `-ufactor=5` (0.5% imbalance tolerance), `-niter=10` (10 refinement iterations), `-seed=42` (reproducible random seed)
- Edge weight handling: Symmetric weights (if edge (u, v) has weight w , then edge (v, u) has weight w)

Hardware Specifications:

- CPU: AMD Ryzen 9 5950X (16 cores, 3.4 GHz base / 4.9 GHz boost)
- RAM: 64 GB DDR4-3600
- Storage: 2 TB NVMe SSD (Samsung 980 PRO)
- OS: Windows 11 Pro (Build 26200), Python 3.13

Runtime Environment:

- Single-threaded execution (METIS does not use parallelism by default)
- No other compute-intensive processes during experiments
- Total wall-clock time for all 43 runs: 47 minutes (mean: 66 seconds per run)

B.2 Data Preprocessing

Census Tract Data:

- Source: US Census Bureau 2020 Redistricting Data (P.L. 94-171)
- Downloaded: April 2021 (official release)
- Tract count by state: Alabama (1,181), Georgia (1,969), Louisiana (1,148), Mississippi (667), South Carolina (1,103)
- Total tracts: 6,068

Minority VAP Calculation:

- Voting Age Population (VAP) = Population aged 18+

- Minority VAP = Black VAP + Hispanic VAP + Asian VAP + Native American VAP + Pacific Islander VAP + Two or More Races VAP
- Non-minority (White) VAP = Total VAP - Minority VAP
- Minority percentage = Minority VAP / Total VAP

Adjacency Graph Construction:

- Method: Shared boundary detection using GeoPandas 0.12.2
- Two tracts are adjacent if their geometries share a boundary segment (not just a point)
- Adjacency is symmetric: if tract i is adjacent o tract j , then tract j is adjacent o tract i
- Average degree: Alabama: 5.8, Georgia: 5.6, Louisiana: 5.9, Mississippi: 5.7, South Carolina: 5.7
- Graph connectivity: All adjacency graphs are connected (single component)

Edge Weight Assignment:

- Minority threshold: $\tau = 0.40$ (40% minority VAP)
- Edge weight factor: $\alpha = 5.0$
- Weight formula: $w(u, v) = \alpha$ if $m_u \geq \tau$ and $m_v \geq \tau$, otherwise $w(u, v) = 1$
- Weight symmetry: $w(u, v) = w(v, u)$ for all edges (u, v)
- Minority-minority edge fraction: Alabama: 12.3%, Georgia: 18.7%, Louisiana: 15.1%, Mississippi: 14.8%, South Carolina: 11.9%

B.3 Experimental Design

Test States: Five states selected based on:

1. High minority population (Black+Hispanic VAP > 35%)
2. Active VRA Section 2 litigation history
3. Range of district counts ($k \in \{4, 6, 7, 14\}$) o test scalability
4. Geographic diversity (Deep South + Atlantic coast)

Partitioning Methods:

1. **Predetermined Recursive Bisection:** For each k , test all possible first split configurations: $(1, k-1), (2, k-2), \dots, (\lfloor k/2 \rfloor, \lceil k/2 \rceil)$. Example: $k = 7$ tests 3 configurations: $(1, 6), (2, 5), (3, 4)$.

2. **Adaptive Recursive Bisection:** Same as predetermined but evaluates all k candidate first splits and selects the one maximizing minority concentration. (In our experiments, adaptive always selected $(1, k - 1)$ due o edge-weighting.)
3. **N-way Partitioning:** Single METIS call with k partitions directly (no recursive splitting).

Total Runs:

- Alabama ($k = 7$): 3 predetermined + 1 adaptive + 1 n-way = 5 runs
- Georgia ($k = 14$): 7 predetermined + 1 adaptive + 1 n-way = 9 runs
- Louisiana ($k = 6$): 3 predetermined + 1 adaptive + 1 n-way = 5 runs
- Mississippi ($k = 4$): 2 predetermined + 1 adaptive + 1 n-way = 4 runs
- South Carolina ($k = 7$): 3 predetermined + 1 adaptive + 1 n-way = 5 runs
- **Total:** 28 runs

Additional Experiments (Section 6.2):

- Robustness test: 10 random seeds for each method imes each state = 150 runs
- Tree structure experiment: All Catalan number trees for $k \in \{4, 5, 6\}$ = 65 additional runs
- **Grand Total:** 243 runs

B.4 Randomness and Reproducibility

METIS Random Seed:

- Fixed seed: `-seed=42` for all runs in main experiment (Section 5)
- Multiple seeds: `-seed=1,2,3,...,10`
for robustness test (Section 6.2)
- Seed affects: tie-breaking during coarsening, initial partition during refinement
- Seed does *not* affect: graph structure, edge weights, balance constraints

Single Run Policy: For the main experiment (Table 2-3, Figures 1-6), we report results from a *single run* per method per state using `-seed=42`. This is justified because:

1. Edge-weighting with $\alpha = 5$ induces deterministic behavior (Theorem 1)
2. Robustness test (Section 6.2) shows variance across seeds is negligible ($< 0.1\%$)
3. Single-run policy matches prior graph partitioning literature ??

For the robustness test, we average over 10 seeds and report mean pm standard deviation.

B.5 Metrics Calculation

Primary Metrics:

- **Max Minority Percentage:** $\max_{i \in [k]} \frac{M_i}{N_i}$ where M_i is minority VAP in district i and N_i is total VAP in district i
- **MM District Count:** Number of districts with minority percentage $\geq 50\%$ (majority-minority districts under VRA)
- **Population Balance:** Maximum deviation from target population. Constraint: $\leq 0.5\%$ (METIS `-ufactor=5`)

Secondary Metrics (Section 5.3):

- **Polsby-Popper Score:** $PP(D) = \frac{4\pi \cdot \text{Area}(D)}{\text{Perimeter}(D)^2}$, computed using GeoPandas `.area` and `.length` methods on dissolved tract geometries
- **Objective Value:** Weighted edge-cut = $\sum_{(u,v) \in \text{cut edges}} w(u,v)$, where cut edge means u and v are in different districts
- **Runtime:** Wall-clock time from start of METIS call to return, measured using Python `time.perf_counter()`

B.6 Statistical Analysis

Variance Calculation: For each state, compute variance across methods:

$$\text{Var}(X) = \frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2$$

where n is number of methods (5-9 depending on state) and x_i is max minority percentage for method i .

Significance Testing: We do not perform statistical significance tests (e.g., t-tests, ANOVA) because:

1. Sample size is small ($n \leq 9$ methods)
2. Observed variance is essentially zero ($< 0.1\%$), making significance testing unnecessary
3. Determinism is theoretical (Theorem 1), not empirical statistical inference

Reporting Precision:

- Minority percentages: 1 decimal place (e.g., 66.2%)
- Compactness scores: 2 decimal places (e.g., 0.42)
- Runtimes: integer seconds (e.g., 83 seconds)
- Variances: 4 decimal places (e.g., 0.0001)

B.7 Code Availability

All code, data, and results are available at:

- GitHub repository: [https://github.com/\[anonymized-for-review\]](https://github.com/[anonymized-for-review])
- Includes: Python scripts, METIS wrapper, tract data processing, visualization code
- License: MIT (open source)
- Documentation: README with step-by-step reproduction instructions

Reproduction Steps:

1. Download census tract shapefiles and redistricting data
2. Run `build_adjacency.py` o construct graphs
3. Run `compute_edge_weights.py` with $\alpha = 5$, $\tau = 0.40$
4. Run `run_partitioning.py` for each method and state
5. Run `analyze_results.py` o compute metrics
6. Run `create_visualizations.py` o generate figures

Expected runtime: < 2 hours on modern hardware.